

Generating SQL for SQLite using Ollama, ChromaDB

This notebook runs through the process of using the `vanna` Python package to generate SQL using AI (RAG + LLMs) including connecting to a database and training. If you're not ready to train on your own database, you can still try it using a sample [SQLite database](#).

Which LLM do you want to use?

- [OpenAI via Vanna.AI \(Recommended\)](#)
Use Vanna.AI for free to generate your queries
- [OpenAI](#)
Use OpenAI with your own API key
- [Azure OpenAI](#)
If you have OpenAI models deployed on Azure
- [\[Selected\] Ollama](#)
Use Ollama locally for free. Requires additional setup.
- [Mistral via Mistral API](#)
If you have a Mistral API key
- [Other LLM](#)
If you have a different LLM model

Where do you want to store the 'training' data?

- [Vanna Hosted Vector DB \(Recommended\)](#)
Use Vanna.AI's hosted vector database (pgvector) for free. This is usable across machines with no additional setup.
- [\[Selected\] ChromaDB](#)
Use ChromaDB's open-source vector database for free locally. No additional setup is necessary -- all database files will be created and stored locally.
- [Marqo](#)
Use Marqo locally for free. Requires additional setup. Or use their hosted option.
- [Other VectorDB](#)

Use any other vector database. Requires additional setup.

Setup

!pip install 'vanna[chromadb]'

```
In [1]: model_name = 'gpt-4o-mini'
        file_db = "~/Downloads/chinook.sqlite"
```

```
In [2]: from api_key_store import ApiKeyStore
        s = ApiKeyStore()

        openai_api_key = s.get_api_key(provider="OPENAI")
```

openai_api_key

```
In [3]: from vanna.openai import OpenAI_Chat
        from vanna.chromadb.chromadb_vector import ChromaDB_VectorStore
```

```
In [4]: class MyVanna(ChromaDB_VectorStore, OpenAI_Chat):
        def __init__(self, config=None):
            ChromaDB_VectorStore.__init__(self, config=config)
            OpenAI_Chat.__init__(self, config=config)

        config = {
            'api_key': openai_api_key,
            'model': model_name
        }
        vn = MyVanna(config=config)
```

Which database do you want to query?

- [Postgres](#)
- [Microsoft SQL Server](#)
- [DuckDB](#)
- [Snowflake](#)
- [BigQuery](#)

- [Selected] SQLite
- [Other Database](#)

[Use Vanna to generate queries for any SQL database](#)

```
In [5]: import os
import re
from time import time
```

```
In [6]: # file_db = "./db/gpt3sql.sqlite"

file_db = os.path.abspath(os.path.expanduser(file_db))
vn.connect_to_sqlite(file_db)
```

```
In [7]: vn.run_sql_is_set
```

```
Out[7]: True
```

```
In [8]: clean_and_train = True # False
```

```
In [9]: hostname = os.uname().nodename
print("Hostname:", hostname)
```

```
Hostname: ducklover1
```

```
In [10]: def remove_collections(collection_name=None, ACCEPTED_TYPES = ["sql", "ddl", "documentation"]):
    if not collection_name:
        collections = ACCEPTED_TYPES
    elif isinstance(collection_name, str):
        collections = [collection_name]
    elif isinstance(collection_name, list):
        collections = collection_name
    else:
        print(f"\t{collection_name} is unknown: Skipped")
        return

    for c in collections:
        if not c in ACCEPTED_TYPES:
            print(f"\t{c} is unknown: Skipped")
```

continue

```
# print(f"vn.remove_collection('{c}')"")
vn.remove_collection(c)
```

```
In [11]: def strip_brackets(ddl):
        """
        This function removes square brackets from table and column names in a DDL script.

        Args:
            ddl (str): The DDL script containing square brackets.

        Returns:
            str: The DDL script with square brackets removed.
        """
        # Use regular expressions to match and replace square brackets
        pattern = r"\[([^\]]+)\]" # Match any character except ] within square brackets
        return re.sub(pattern, r"\1", ddl)
```

```
In [12]: if clean_and_train:
        remove_collections()
```

Training

You only need to train once. Do not train again unless you want to add more training data.

```
In [13]: # show training data
training_data = vn.get_training_data()
training_data
```

```
Out[13]: id question content training_data_type
```

```
In [14]: df_ddl = vn.run_sql("SELECT type, sql FROM sqlite_master WHERE sql is not null")
```

```
In [15]: df_ddl
```

Out[15]:

	type	sql
0	table	CREATE TABLE "albums"\r\n(\r\n [AlbumId] IN...
1	table	CREATE TABLE sqlite_sequence(name,seq)
2	table	CREATE TABLE "artists"\r\n(\r\n [ArtistId] ...
3	table	CREATE TABLE "customers"\r\n(\r\n [Customer...
4	table	CREATE TABLE "employees"\r\n(\r\n [Employee...
5	table	CREATE TABLE "genres"\r\n(\r\n [GenreId] IN...
6	table	CREATE TABLE "invoices"\r\n(\r\n [InvoiceId]...
7	table	CREATE TABLE "invoice_items"\r\n(\r\n [Invo...
8	table	CREATE TABLE "media_types"\r\n(\r\n [MediaT...
9	table	CREATE TABLE "playlists"\r\n(\r\n [Playlist...
10	table	CREATE TABLE "playlist_track"\r\n(\r\n [Pla...
11	table	CREATE TABLE "tracks"\r\n(\r\n [TrackId] IN...
12	index	CREATE INDEX [IFK_AlbumArtistId] ON "albums" (...
13	index	CREATE INDEX [IFK_CustomerSupportRepId] ON "cu...
14	index	CREATE INDEX [IFK_EmployeeReportsTo] ON "emplo...
15	index	CREATE INDEX [IFK_InvoiceCustomerId] ON "invoi...
16	index	CREATE INDEX [IFK_InvoiceLineInvoiceId] ON "in...
17	index	CREATE INDEX [IFK_InvoiceLineTrackId] ON "invo...
18	index	CREATE INDEX [IFK_PlaylistTrackTrackId] ON "pl...
19	index	CREATE INDEX [IFK_TrackAlbumId] ON "tracks" ([...
20	index	CREATE INDEX [IFK_TrackGenreId] ON "tracks" ([...
21	index	CREATE INDEX [IFK_TrackMediaTypeId] ON "tracks...
22	table	CREATE TABLE sqlite_stat1(tbl,idx,stat)

```
In [16]: if clean_and_train:
        for ddl in df_ddl['sql'].to_list():
            ddl = strip_brackets(ddl)
            vn.train(ddl=ddl)

        # Sometimes you may want to add documentation about your business terminology or definitions.
        vn.train(documentation="In the chinook database invoice means order")
```

Adding ddl: CREATE TABLE "albums"

```
(
  AlbumId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
  Title NVARCHAR(160) NOT NULL,
  ArtistId INTEGER NOT NULL,
  FOREIGN KEY (ArtistId) REFERENCES "artists" (ArtistId)
    ON DELETE NO ACTION ON UPDATE NO ACTION
)
```

Adding ddl: CREATE TABLE sqlite_sequence(name,seq)

Adding ddl: CREATE TABLE "artists"

```
(
  ArtistId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
  Name NVARCHAR(120)
)
```

Adding ddl: CREATE TABLE "customers"

```
(
  CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
  FirstName NVARCHAR(40) NOT NULL,
  LastName NVARCHAR(20) NOT NULL,
  Company NVARCHAR(80),
  Address NVARCHAR(70),
  City NVARCHAR(40),
  State NVARCHAR(40),
  Country NVARCHAR(40),
  PostalCode NVARCHAR(10),
  Phone NVARCHAR(24),
  Fax NVARCHAR(24),
  Email NVARCHAR(60) NOT NULL,
  SupportRepId INTEGER,
  FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId)
    ON DELETE NO ACTION ON UPDATE NO ACTION
)
```

Adding ddl: CREATE TABLE "employees"

```
(
  EmployeeId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
  LastName NVARCHAR(20) NOT NULL,
  FirstName NVARCHAR(20) NOT NULL,
  Title NVARCHAR(30),
  ReportsTo INTEGER,
  BirthDate DATETIME,
  HireDate DATETIME,
)
```

```
Address NVARCHAR(70),
City NVARCHAR(40),
State NVARCHAR(40),
Country NVARCHAR(40),
PostalCode NVARCHAR(10),
Phone NVARCHAR(24),
Fax NVARCHAR(24),
Email NVARCHAR(60),
FOREIGN KEY (ReportsTo) REFERENCES "employees" (EmployeeId)
    ON DELETE NO ACTION ON UPDATE NO ACTION
)
Adding ddl: CREATE TABLE "genres"
(
    GenreId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
    Name NVARCHAR(120)
)
Adding ddl: CREATE TABLE "invoices"
(
    InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
    CustomerId INTEGER NOT NULL,
    InvoiceDate DATETIME NOT NULL,
    BillingAddress NVARCHAR(70),
    BillingCity NVARCHAR(40),
    BillingState NVARCHAR(40),
    BillingCountry NVARCHAR(40),
    BillingPostalCode NVARCHAR(10),
    Total NUMERIC(10,2) NOT NULL,
    FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId)
        ON DELETE NO ACTION ON UPDATE NO ACTION
)
Adding ddl: CREATE TABLE "invoice_items"
(
    InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
    InvoiceId INTEGER NOT NULL,
    TrackId INTEGER NOT NULL,
    UnitPrice NUMERIC(10,2) NOT NULL,
    Quantity INTEGER NOT NULL,
    FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId)
        ON DELETE NO ACTION ON UPDATE NO ACTION,
    FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId)
        ON DELETE NO ACTION ON UPDATE NO ACTION
)
```



```
)
Adding ddl: CREATE TABLE "media_types"
(
    MediaTypeId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
    Name NVARCHAR(120)
)
Adding ddl: CREATE TABLE "playlists"
(
    PlaylistId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
    Name NVARCHAR(120)
)
Adding ddl: CREATE TABLE "playlist_track"
(
    PlaylistId INTEGER NOT NULL,
    TrackId INTEGER NOT NULL,
    CONSTRAINT PK_PlaylistTrack PRIMARY KEY (PlaylistId, TrackId),
    FOREIGN KEY (PlaylistId) REFERENCES "playlists" (PlaylistId)
        ON DELETE NO ACTION ON UPDATE NO ACTION,
    FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId)
        ON DELETE NO ACTION ON UPDATE NO ACTION
)
Adding ddl: CREATE TABLE "tracks"
(
    TrackId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
    Name NVARCHAR(200) NOT NULL,
    AlbumId INTEGER,
    MediaTypeId INTEGER NOT NULL,
    GenreId INTEGER,
    Composer NVARCHAR(220),
    Milliseconds INTEGER NOT NULL,
    Bytes INTEGER,
    UnitPrice NUMERIC(10,2) NOT NULL,
    FOREIGN KEY (AlbumId) REFERENCES "albums" (AlbumId)
        ON DELETE NO ACTION ON UPDATE NO ACTION,
    FOREIGN KEY (GenreId) REFERENCES "genres" (GenreId)
        ON DELETE NO ACTION ON UPDATE NO ACTION,
    FOREIGN KEY (MediaTypeId) REFERENCES "media_types" (MediaTypeId)
        ON DELETE NO ACTION ON UPDATE NO ACTION
)
Adding ddl: CREATE INDEX IFK_AlbumArtistId ON "albums" (ArtistId)
Adding ddl: CREATE INDEX IFK_CustomerSupportRepId ON "customers" (SupportRepId)
```

```
Adding ddl: CREATE INDEX IFK_EmployeeReportsTo ON "employees" (ReportsTo)
Adding ddl: CREATE INDEX IFK_InvoiceCustomerId ON "invoices" (CustomerId)
Adding ddl: CREATE INDEX IFK_InvoiceLineInvoiceId ON "invoice_items" (InvoiceId)
Adding ddl: CREATE INDEX IFK_InvoiceLineTrackId ON "invoice_items" (TrackId)
Adding ddl: CREATE INDEX IFK_PlaylistTrackTrackId ON "playlist_track" (TrackId)
Adding ddl: CREATE INDEX IFK_TrackAlbumId ON "tracks" (AlbumId)
Adding ddl: CREATE INDEX IFK_TrackGenreId ON "tracks" (GenreId)
Adding ddl: CREATE INDEX IFK_TrackMediaTypeId ON "tracks" (MediaTypeId)
Adding ddl: CREATE TABLE sqlite_stat1(tbl,idx,stat)
Adding documentation....
```

In []:

Asking the AI

Whenever you ask a new question, it will find the 10 most relevant pieces of training data and use it as part of the LLM prompt to generate the SQL.

In [17]: `ts_start = time()`

In [18]: `vn.ask(question="Show me a list of tables in the SQLite database")`

Number of requested results 10 is greater than number of elements in index 1, updating n_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables \nCREATE TABLE sqlite_stat1(tbl,idx,stat)\n\nCREATE TABLE sqlite_sequence(name,seq)\n\nCREATE TABLE "playlists"\n\n(\n PlaylistId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(120)\n)\n\nCREATE TABLE "media_types"\n\n(\n MediaTypeId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(120)\n)\n\nCREATE TABLE "artists"\n\n(\n ArtistId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(120)\n)\n\nCREATE TABLE "genres"\n\n(\n GenreId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(120)\n)\n\nCREATE TABLE "invoice_items"\n\n(\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n InvoiceId INTEGER NOT NULL,\n TrackId INTEGER NOT NULL,\n UnitPrice NUMERIC(10,2) NOT NULL,\n Quantity INTEGER NOT NULL,\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n)\n\nCREATE TABLE "tracks"\n\n(\n TrackId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(200) NOT NULL,\n AlbumId INTEGER,\n MediaTypeId INTEGER NOT NULL,\n GenreId INTEGER,\n Composer NVARCHAR(220),\n Milliseconds INTEGER NOT NULL,\n Bytes INTEGER,\n UnitPrice NUMERIC(10,2) NOT NULL,\n FOREIGN KEY (AlbumId) REFERENCES "albums" (AlbumId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (GenreId) REFERENCES "genres" (GenreId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (MediaTypeId) REFERENCES "media_types" (MediaTypeId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n)\n\n\n===Additional Context \n\nIn the chinook database invoice means order\n\n===Response Guidelines \n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question. \n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate_sql \n3. If the provided context is insufficient, please explain why it can't be generated. \n4. Please use the most relevant table(s). \n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before. \n'}], {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}]

Using model gpt-4o-mini for 652.0 tokens (approx)

LLM Response: ```sql

```
SELECT name FROM sqlite_master WHERE type='table';
```
```

Extracted SQL: SELECT name FROM sqlite\_master WHERE type='table';

```
SELECT name FROM sqlite_master WHERE type='table';
```

|   | name            |
|---|-----------------|
| 0 | albums          |
| 1 | sqlite_sequence |
| 2 | artists         |
| 3 | customers       |
| 4 | employees       |
| 5 | genres          |
| 6 | invoices        |
| 7 | invoice_items   |
| 8 | media_types     |


9        playlists

10     playlist\_track

11        tracks

12     sqlite\_stat1

Using model gpt-4o-mini for 168.0 tokens (approx)



```

Out[18]: ("SELECT name FROM sqlite_master WHERE type='table';",
 name
0 albums
1 sqlite_sequence
2 artists
3 customers
4 employees
5 genres
6 invoices
7 invoice_items
8 media_types
9 playlists
10 playlist_track
11 tracks
12 sqlite_stat1,
Figure({
 'data': [{'type': 'bar',
 'x': array(['albums', 'sqlite_sequence', 'artists', 'customers', 'employees',
 'genres', 'invoices', 'invoice_items', 'media_types', 'playlists',
 'playlist_track', 'tracks', 'sqlite_stat1'], dtype=object),
 'y': [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]}],
 'layout': {'template': '...'}}))

```

```
In [19]: vn.ask(question="which table stores customer's orders")
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1  
 Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n CustomerId INTEGER NOT NULL,\n InvoiceDate DATETIME NOT NULL,\n BillingAddress NVARCHAR(70),\n BillingCity NVARCHAR(40),\n BillingState NVARCHAR(40),\n BillingCountry NVARCHAR(40),\n BillingPostalCode NVARCHAR(10),\n Total NUMERIC(10,2) NOT NULL,\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId) \n\n ON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n InvoiceId INTEGER NOT NULL,\n TrackId INTEGER NOT NULL,\n UnitPrice NUMERIC(10,2) NOT NULL,\n Quantity INTEGER NOT NULL,\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\n ON DELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\n ON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE TABLE "customers"\n\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n FirstName NVARCHAR(40) NOT NULL,\n LastName NVARCHAR(20) NOT NULL,\n Company NVARCHAR(80),\n Address NVARCHAR(70),\n City NVARCHAR(40),\n State NVARCHAR(40),\n Country NVARCHAR(40),\n PostalCode NVARCHAR(10),\n Phone NVARCHAR(24),\n Fax NVARCHAR(24),\n Email NVARCHAR(60) NOT NULL,\n SupportRepId INTEGER,\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId) \n\n ON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE TABLE sqlite\_sequence (name,seq)\n\nCREATE TABLE "playlists"\n\n PlaylistId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(120)\n\nCREATE TABLE sqlite\_stat1(tbl,idx,stat)\n\n===Additional Context\n\nIn the chinook database invoice means order\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'}], {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': "which table stores customer's orders"}]

Using model gpt-4o-mini for 683.75 tokens (approx)

LLM Response: The "invoices" table stores customer's orders.

The "invoices" table stores customer's orders.

Couldn't run sql: Execution failed on sql 'The "invoices" table stores customer's orders.': near "The": syntax error

In [20]: `vn.ask(question="How many records are in table called customer")`

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1  
 Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n CustomerId INTEGER NOT NULL,\n\n InvoiceDate DATETIME NOT NULL,\n\n BillingAddress NVARCHAR(70),\n\n BillingCity NVARCHAR(40),\n\n BillingState NVARCHAR(40),\n\n BillingCountry NVARCHAR(40),\n\n BillingPostalCode NVARCHAR(10),\n\n Total NUMERIC(10,2) NOT NULL,\n\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE TABLE "customers"\n\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n FirstName NVARCHAR(40) NOT NULL,\n\n LastName NVARCHAR(20) NOT NULL,\n\n Company NVARCHAR(80),\n\n Address NVARCHAR(70),\n\n City NVARCHAR(40),\n\n State NVARCHAR(40),\n\n Country NVARCHAR(40),\n\n PostalCode NVARCHAR(10),\n\n Phone NVARCHAR(24),\n\n Fax NVARCHAR(24),\n\n Email NVARCHAR(60) NOT NULL,\n\n SupportRepId INTEGER,\n\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n InvoiceId INTEGER NOT NULL,\n\n TrackId INTEGER NOT NULL,\n\n UnitPrice NUMERIC(10,2) NOT NULL,\n\n Quantity INTEGER NOT NULL,\n\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE TABLE "albums"\n\n AlbumId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n Title NVARCHAR(160) NOT NULL,\n\n ArtistId INTEGER NOT NULL,\n\n FOREIGN KEY (ArtistId) REFERENCES "artists" (ArtistId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\n\n===Additional Context\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'}], {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': 'How many records are in table called customer'}]

Using model gpt-4o-mini for 690.25 tokens (approx)

LLM Response: SELECT COUNT(\*) FROM customers;

Extracted SQL: SELECT COUNT(\*) FROM customers;

SELECT COUNT(\*) FROM customers;

COUNT(\*)

0 59

Using model gpt-4o-mini for 163.5 tokens (approx)



Total Customer Records

59

```
Out[20]: ('SELECT COUNT(*) FROM customers;',
 COUNT(*)
 0 59,
 Figure({
 'data': [{'mode': 'number', 'title': {'text': 'Total Customer Records'}, 'type': 'indicator', 'value': 59}],
 'layout': {'template': '...'}
 })))
```

```
In [21]: vn.ask(question="How many customers are there")
```

Number of requested results 10 is greater than number of elements in index 2, updating n\_results = 2

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{ 'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n CustomerId INTEGER NOT NULL,\n\n InvoiceDate DATETIME NOT NULL,\n\n BillingAddress NVARCHAR(70),\n\n BillingCity NVARCHAR(40),\n\n BillingState NVARCHAR(40),\n\n BillingCountry NVARCHAR(40),\n\n BillingPostalCode NVARCHAR(10),\n\n Total NUMERIC(10,2) NOT NULL,\n\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE INDEX IFK\_CustomerSupportRepId ON "customers" (SupportRepId)\n\n\nCREATE TABLE "customers"\n\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n FirstName NVARCHAR(40) NOT NULL,\n\n LastName NVARCHAR(20) NOT NULL,\n\n Company NVARCHAR(80),\n\n Address NVARCHAR(70),\n\n City NVARCHAR(40),\n\n State NVARCHAR(40),\n\n Country NVARCHAR(40),\n\n PostalCode NVARCHAR(10),\n\n Phone NVARCHAR(24),\n\n Fax NVARCHAR(24),\n\n Email NVARCHAR(60) NOT NULL,\n\n SupportRepId INTEGER,\n\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE INDEX IFK\_InvoiceCustomerId ON "invoices" (CustomerId)\n\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n InvoiceId INTEGER NOT NULL,\n\n TrackId INTEGER NOT NULL,\n\n UnitPrice NUMERIC(10,2) NOT NULL,\n\n Quantity INTEGER NOT NULL,\n\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE INDEX IFK\_InvoiceLineInvoiceId ON "invoice\_items" (InvoiceId)\n\n\nCREATE INDEX IFK\_InvoiceLineTrackId ON "invoice\_items" (TrackId)\n\n\n\n===Additional Context\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'], { 'role': 'user', 'content': 'How many records are in table called customer'}, { 'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, { 'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, { 'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, { 'role': 'user', 'content': 'How many customers are there'}]

Using model gpt-4o-mini for 707.25 tokens (approx)

LLM Response: SELECT COUNT(\*) FROM customers;

Extracted SQL: SELECT COUNT(\*) FROM customers;

SELECT COUNT(\*) FROM customers;

COUNT(\*)

0 59

Using model gpt-4o-mini for 159.25 tokens (approx)

Total Customers

59

```
Out[21]: ('SELECT COUNT(*) FROM customers;',
 COUNT(*)
 0 59,
 Figure({
 'data': [{'mode': 'number', 'title': {'text': 'Total Customers'}, 'type': 'indicator', 'value': 59}],
 'layout': {'template': '...'}
 })))
```

In [ ]:

```
In [22]: vn.ask(question="what are the top 5 countries that customers come from?")
```

Number of requested results 10 is greater than number of elements in index 3, updating n\_results = 3  
Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

```
Couldn't run sql: Execution failed on sql 'The LLM is not allowed to see the data in your database. Your question requires database introspection to generate the necessary SQL. Please set allow_llm_to_see_data=True to enable this feature.': near "The": syntax error
```

see `sample-sql-queries-sqlite-chinook.ipynb`

```
In [23]: question = """
 List all albums and their corresponding artist names
 """

vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 3, updating n\_results = 3

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

1 Balls to the Wall

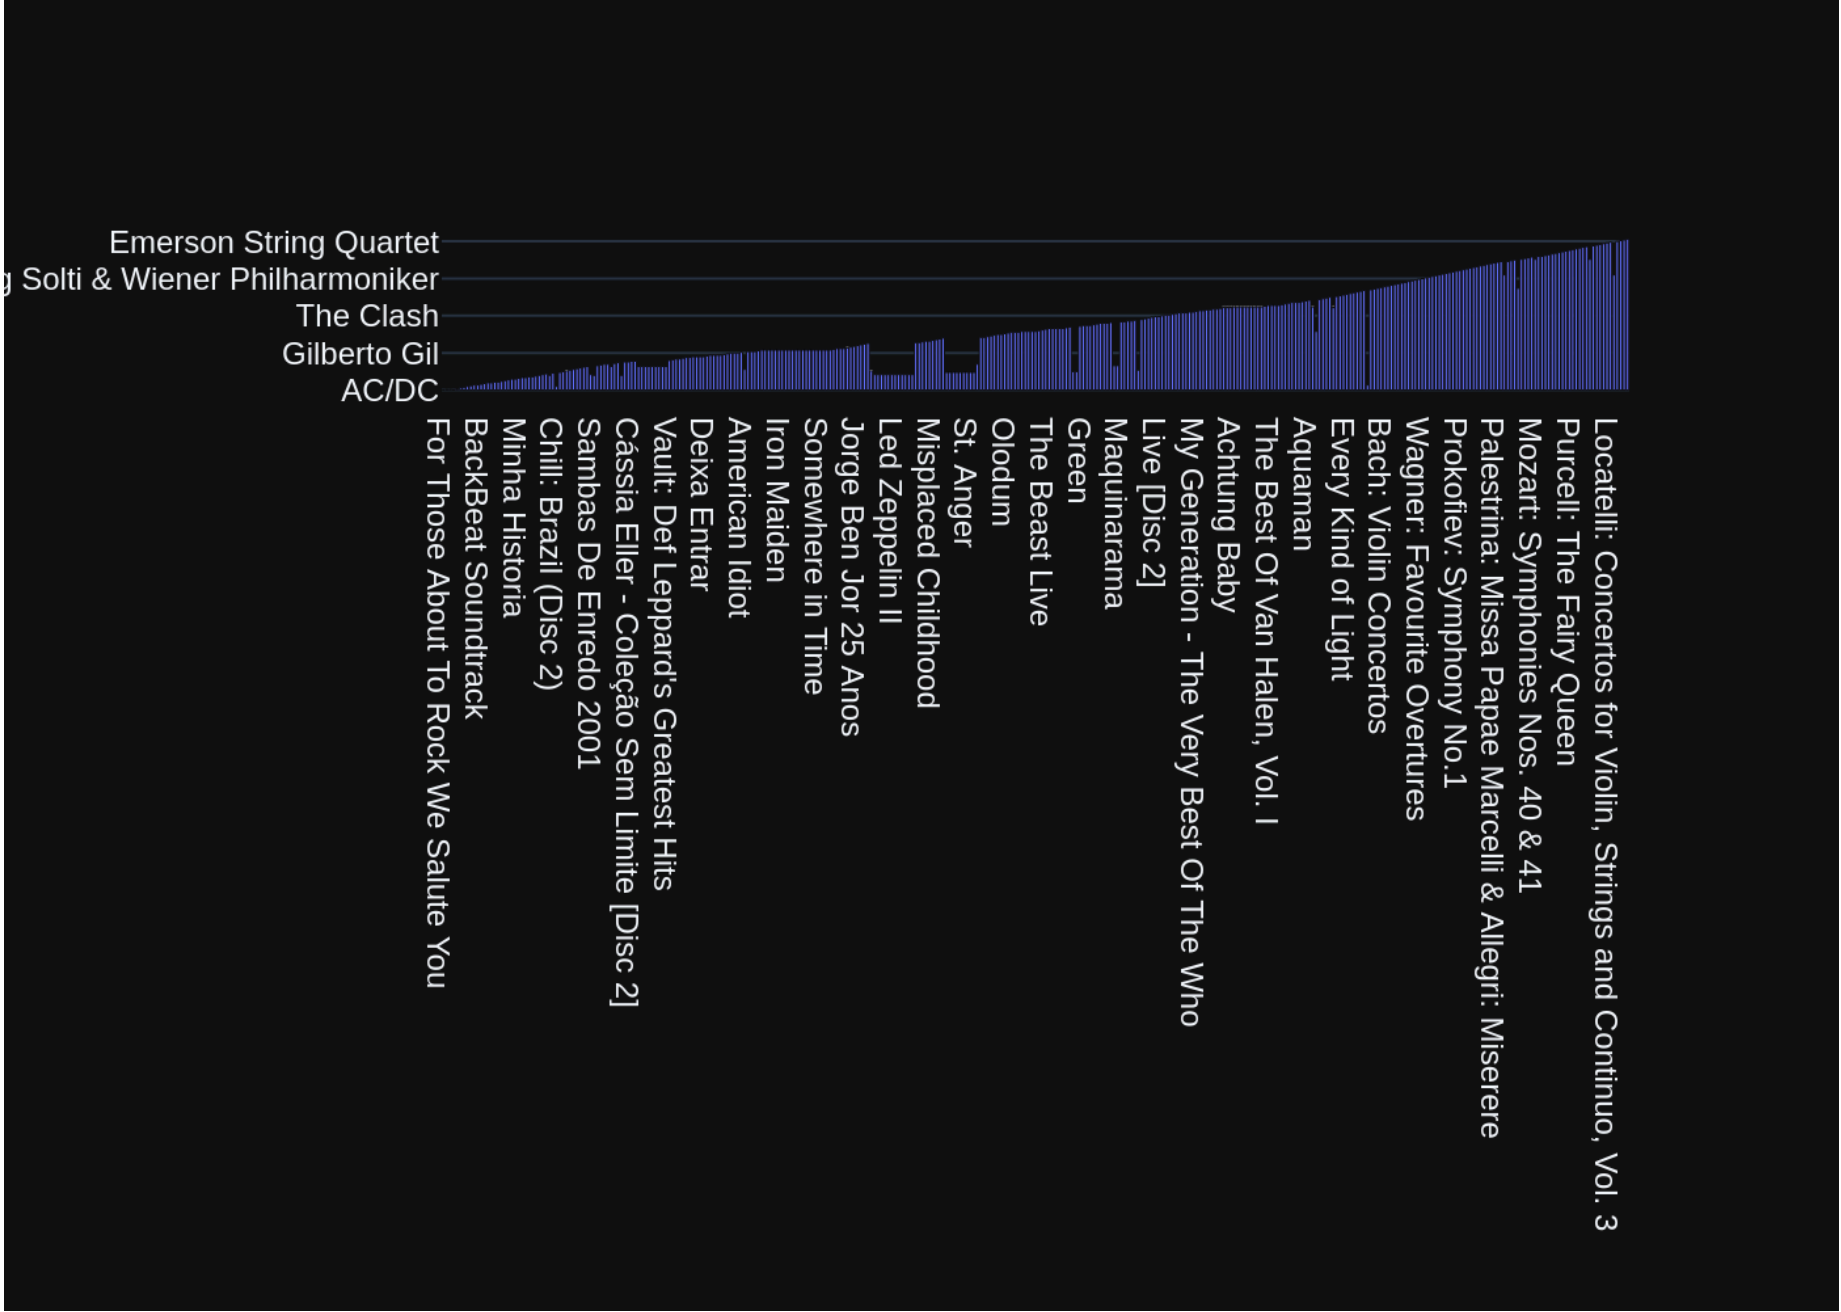


```
2 Restless and Wild
3 Let There Be Rock
4 Big Ones
..
342 Respighi:Pines of Rome
343 Schubert: The Late String Quartets & String Qu...
344 Monteverdi: L'Orfeo
345 Mozart: Chamber Music
346 Koyaanisqatsi (Soundtrack from the Motion Pict...
```

```
ArtistName
0 AC/DC
1 Accept
2 Accept
3 AC/DC
4 Aerosmith
..
342 Eugene Ormandy
343 Emerson String Quartet
344 C. Monteverdi, Nigel Rogers - Chiaroscuro; Lon...
345 Nash Ensemble
346 Philip Glass Ensemble
```

[347 rows x 2 columns]

Using model gpt-4o-mini for 197.75 tokens (approx)



```
Out[23]: ('SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId =
artists.ArtistId;',
```

```

AlbumTitle \
0 For Those About To Rock We Salute You
1 Balls to the Wall
2 Restless and Wild
3 Let There Be Rock
4 Big Ones
```

```

.. ...
342 Respighi:Pines of Rome
343 Schubert: The Late String Quartets & String Qu...
344 Monteverdi: L'Orfeo
345 Mozart: Chamber Music
346 Koyaanisqatsi (Soundtrack from the Motion Pict...
```

```

ArtistName
0 AC/DC
1 Accept
2 Accept
3 AC/DC
4 Aerosmith
```

```

.. ...
342 Eugene Ormandy
343 Emerson String Quartet
344 C. Monteverdi, Nigel Rogers - Chiaroscuro; Lon...
345 Nash Ensemble
346 Philip Glass Ensemble
```

```
[347 rows x 2 columns],
```

```
Figure({
 'data': [{'text': array(['AC/DC', 'Accept', 'Accept', ...,
 'C. Monteverdi, Nigel Rogers - Chiaroscuro; London Baroque; London Cornett & Sackb
u',
 'Nash Ensemble', 'Philip Glass Ensemble'], dtype=object),
 'textposition': 'auto',
 'type': 'bar',
 'x': array(['For Those About To Rock We Salute You', 'Balls to the Wall',
 'Restless and Wild', ..., 'Monteverdi: L'Orfeo',
 'Mozart: Chamber Music',
 'Koyaanisqatsi (Soundtrack from the Motion Picture)'], dtype=object),
```

```
 'y': array(['AC/DC', 'Accept', 'Accept', ...,
 'C. Monteverdi, Nigel Rogers - Chiaroscuro; London Baroque; London Cornett & Sackbu',
 'Nash Ensemble', 'Philip Glass Ensemble'], dtype=object)]],
 'layout': {'template': '...'}
))
```

```
In [24]: question = """
 Find all tracks with a name containing "What" (case-insensitive)
 """

 vn.ask(question=question)
```

```
Number of requested results 10 is greater than number of elements in index 4, updating n_results = 4
Number of requested results 10 is greater than number of elements in index 1, updating n_results = 1
```

```
LLM Response: SELECT * FROM tracks WHERE Name LIKE '%What%';
Extracted SQL: SELECT * FROM tracks WHERE Name LIKE '%What%';
SELECT * FROM tracks WHERE Name LIKE '%What%';
```


|   | TrackId | Name                        | AlbumId |
|---|---------|-----------------------------|---------|
| 0 | 26      | What It Takes               | 5       |
| 1 | 88      | What You Are                | 10      |
| 2 | 130     | Do what cha wanna           | 13      |
| 3 | 342     | What is and Should Never Be | 30      |
| 4 | 607     | So What                     | 48      |

|    |      |                                              |     |
|----|------|----------------------------------------------|-----|
| 5  | 960  | What A Day                                   | 76  |
| 6  | 1000 | What If I Do?                                | 80  |
| 7  | 1039 | What Now My Love                             | 83  |
| 8  | 1145 | Whatsername                                  | 89  |
| 9  | 1440 | Whatever It Is, I Just Can't Stop            | 116 |
| 10 | 1469 | Look What You've Done                        | 119 |
| 11 | 1470 | Get What You Need                            | 119 |
| 12 | 1628 | What Is And What Should Never Be             | 133 |
| 13 | 1778 | You're What's Happening (In The World Today) | 146 |
| 14 | 1823 | So What                                      | 149 |
| 15 | 2772 | I Don't Know What To Do With Myself          | 223 |
| 16 | 2884 | What Kate Did                                | 231 |
| 17 | 2893 | Whatever the Case May Be                     | 230 |
| 18 | 2992 | I Still Haven't Found What I'm Looking for   | 237 |
| 19 | 3007 | I Still Haven't Found What I'm Looking For   | 238 |
| 20 | 3258 | Whatever Gets You Thru the Night             | 255 |
| 21 | 3475 | What Is It About Men                         | 322 |

|    | MediaTypeId | GenreId | Composer \                                        |
|----|-------------|---------|---------------------------------------------------|
| 0  | 1           | 1       | Steven Tyler, Joe Perry, Desmond Child            |
| 1  | 1           | 1       | Audioslave/Chris Cornell                          |
| 2  | 1           | 2       | George Duke                                       |
| 3  | 1           | 1       | Jimmy Page/Robert Plant                           |
| 4  | 1           | 2       | Miles Davis                                       |
| 5  | 1           | 1       | Mike Bordin, Billy Gould, Mike Patton             |
| 6  | 1           | 1       | Dave Grohl, Taylor Hawkins, Nate Mendel, Chris... |
| 7  | 1           | 12      | carl sigman/gilbert becaud/pierre leroyer         |
| 8  | 1           | 4       | Green Day                                         |
| 9  | 1           | 1       | Jay Kay/Kay, Jay                                  |
| 10 | 1           | 4       | N. Cester                                         |
| 11 | 1           | 4       | C. Cester/C. Muncey/N. Cester                     |
| 12 | 1           | 1       | Jimmy Page, Robert Plant                          |
| 13 | 1           | 14      | Allen Story/George Gordy/Robert Gordy             |
| 14 | 1           | 3       | Culmer/Exalt                                      |
| 15 | 1           | 7       | None                                              |
| 16 | 3           | 19      | None                                              |
| 17 | 3           | 19      | None                                              |
| 18 | 1           | 1       | Bono/Clayton, Adam/Mullen Jr., Larry/The Edge     |
| 19 | 1           | 1       | U2                                                |
| 20 | 2           | 9       | None                                              |
| 21 | 2           | 9       | Delroy "Chris" Cooper, Donovan Jackson, Earl C... |

|    | Milliseconds | Bytes     | UnitPrice |
|----|--------------|-----------|-----------|
| 0  | 310622       | 10144730  | 0.99      |
| 1  | 249391       | 5988186   | 0.99      |
| 2  | 274155       | 9018565   | 0.99      |
| 3  | 260675       | 8497116   | 0.99      |
| 4  | 564009       | 18360449  | 0.99      |
| 5  | 158275       | 5203430   | 0.99      |
| 6  | 302994       | 9929799   | 0.99      |
| 7  | 149995       | 4913383   | 0.99      |
| 8  | 252316       | 8244843   | 0.99      |
| 9  | 247222       | 8249453   | 0.99      |
| 10 | 230974       | 7517083   | 0.99      |
| 11 | 247719       | 8043765   | 0.99      |
| 12 | 287973       | 9369385   | 0.99      |
| 13 | 142027       | 4631104   | 0.99      |
| 14 | 189152       | 6162894   | 0.99      |
| 15 | 221387       | 7251478   | 0.99      |
| 16 | 2610250      | 484583988 | 1.99      |
| 17 | 2616410      | 183867185 | 1.99      |
| 18 | 353567       | 11542247  | 0.99      |
| 19 | 280764       | 9306737   | 0.99      |
| 20 | 215084       | 3499018   | 0.99      |
| 21 | 209573       | 3426106   | 0.99      |

Using model gpt-4o-mini for 223.5 tokens (approx)





Out[24]: ("SELECT \* FROM tracks WHERE Name LIKE '%What%';",

|    | TrackId | Name                                         | AlbumId | \ |
|----|---------|----------------------------------------------|---------|---|
| 0  | 26      | What It Takes                                | 5       |   |
| 1  | 88      | What You Are                                 | 10      |   |
| 2  | 130     | Do what cha wanna                            | 13      |   |
| 3  | 342     | What is and Should Never Be                  | 30      |   |
| 4  | 607     | So What                                      | 48      |   |
| 5  | 960     | What A Day                                   | 76      |   |
| 6  | 1000    | What If I Do?                                | 80      |   |
| 7  | 1039    | What Now My Love                             | 83      |   |
| 8  | 1145    | Whatsername                                  | 89      |   |
| 9  | 1440    | Whatever It Is, I Just Can't Stop            | 116     |   |
| 10 | 1469    | Look What You've Done                        | 119     |   |
| 11 | 1470    | Get What You Need                            | 119     |   |
| 12 | 1628    | What Is And What Should Never Be             | 133     |   |
| 13 | 1778    | You're What's Happening (In The World Today) | 146     |   |
| 14 | 1823    | So What                                      | 149     |   |
| 15 | 2772    | I Don't Know What To Do With Myself          | 223     |   |
| 16 | 2884    | What Kate Did                                | 231     |   |
| 17 | 2893    | Whatever the Case May Be                     | 230     |   |
| 18 | 2992    | I Still Haven't Found What I'm Looking for   | 237     |   |
| 19 | 3007    | I Still Haven't Found What I'm Looking For   | 238     |   |
| 20 | 3258    | Whatever Gets You Thru the Night             | 255     |   |
| 21 | 3475    | What Is It About Men                         | 322     |   |

|    | MediaTypeId | GenreId | Composer                                          | \ |
|----|-------------|---------|---------------------------------------------------|---|
| 0  | 1           | 1       | Steven Tyler, Joe Perry, Desmond Child            |   |
| 1  | 1           | 1       | Audioslave/Chris Cornell                          |   |
| 2  | 1           | 2       | George Duke                                       |   |
| 3  | 1           | 1       | Jimmy Page/Robert Plant                           |   |
| 4  | 1           | 2       | Miles Davis                                       |   |
| 5  | 1           | 1       | Mike Bordin, Billy Gould, Mike Patton             |   |
| 6  | 1           | 1       | Dave Grohl, Taylor Hawkins, Nate Mendel, Chris... |   |
| 7  | 1           | 12      | carl sigman/gilbert becaud/pierre leroyer         |   |
| 8  | 1           | 4       | Green Day                                         |   |
| 9  | 1           | 1       | Jay Kay/Kay, Jay                                  |   |
| 10 | 1           | 4       | N. Cester                                         |   |
| 11 | 1           | 4       | C. Cester/C. Muncey/N. Cester                     |   |
| 12 | 1           | 1       | Jimmy Page, Robert Plant                          |   |
| 13 | 1           | 14      | Allen Story/George Gordy/Robert Gordy             |   |

|    |   |    |                                                   |
|----|---|----|---------------------------------------------------|
| 14 | 1 | 3  | Culmer/Exalt                                      |
| 15 | 1 | 7  | None                                              |
| 16 | 3 | 19 | None                                              |
| 17 | 3 | 19 | None                                              |
| 18 | 1 | 1  | Bono/Clayton, Adam/Mullen Jr., Larry/The Edge     |
| 19 | 1 | 1  | U2                                                |
| 20 | 2 | 9  | None                                              |
| 21 | 2 | 9  | Delroy "Chris" Cooper, Donovan Jackson, Earl C... |

|    | Milliseconds | Bytes     | UnitPrice |
|----|--------------|-----------|-----------|
| 0  | 310622       | 10144730  | 0.99      |
| 1  | 249391       | 5988186   | 0.99      |
| 2  | 274155       | 9018565   | 0.99      |
| 3  | 260675       | 8497116   | 0.99      |
| 4  | 564009       | 18360449  | 0.99      |
| 5  | 158275       | 5203430   | 0.99      |
| 6  | 302994       | 9929799   | 0.99      |
| 7  | 149995       | 4913383   | 0.99      |
| 8  | 252316       | 8244843   | 0.99      |
| 9  | 247222       | 8249453   | 0.99      |
| 10 | 230974       | 7517083   | 0.99      |
| 11 | 247719       | 8043765   | 0.99      |
| 12 | 287973       | 9369385   | 0.99      |
| 13 | 142027       | 4631104   | 0.99      |
| 14 | 189152       | 6162894   | 0.99      |
| 15 | 221387       | 7251478   | 0.99      |
| 16 | 2610250      | 484583988 | 1.99      |
| 17 | 2616410      | 183867185 | 1.99      |
| 18 | 353567       | 11542247  | 0.99      |
| 19 | 280764       | 9306737   | 0.99      |
| 20 | 215084       | 3499018   | 0.99      |
| 21 | 209573       | 3426106   | 0.99 ,    |

```
Figure({
 'data': [{'marker': {'color': 'lightsalmon'}},
 'text': array(['Steven Tyler, Joe Perry, Desmond Child', 'Audioslave/Chris Cornell',
 'George Duke', 'Jimmy Page/Robert Plant', 'Miles Davis',
 'Mike Bordin, Billy Gould, Mike Patton',
 'Dave Grohl, Taylor Hawkins, Nate Mendel, Chris Shiflett/FOO FIGHTERS',
 'carl sigman/gilbert becaud/pierre leroyer', 'Green Day',
 'Jay Kay/Kay, Jay', 'N. Cester', 'C. Cester/C. Muncey/N. Cester',
 'Jimmy Page, Robert Plant', 'Allen Story/George Gordy/Robert Gordy',
```

```

 'Culmer/Exalt', None, None, None,
 'Bono/Clayton, Adam/Mullen Jr., Larry/The Edge', 'U2', None,
 'Delroy "Chris" Cooper, Donovan Jackson, Earl Chinna Smith, Felix Howard, Gordon Wil
 liams, Luke Smith, Paul Watson & Wilburn Squiddley Cole'],
 dtype=object),
 'type': 'bar',
 'x': array(['What It Takes', 'What You Are', 'Do what cha wanna',
 'What is and Should Never Be', 'So What', 'What A Day', 'What If I Do?',
 'What Now My Love', 'Whatsername', "Whatever It Is, I Just Can't Stop",
 "Look What You've Done", 'Get What You Need',
 'What Is And What Should Never Be',
 "You're What's Happening (In The World Today)", 'So What',
 "I Don't Know What To Do With Myself", 'What Kate Did',
 'Whatever the Case May Be',
 "I Still Haven't Found What I'm Looking for",
 "I Still Haven't Found What I'm Looking For",
 'Whatever Gets You Thru the Night', 'What Is It About Men'],
 dtype=object),
 'y': array([0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99, 0.99,
 0.99, 0.99, 0.99, 0.99, 1.99, 1.99, 0.99, 0.99, 0.99, 0.99])),
 'layout': {'template': '...',
 'title': {'text': 'Tracks with "What" in Name'},
 'xaxis': {'title': {'text': 'Track Name'}},
 'yaxis': {'title': {'text': 'Unit Price'}}}
)))

```

```

In [25]: question = """
 Get the total number of invoices for each customer
 """

 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 5, updating n\_results = 5  
 Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n CustomerId INTEGER NOT NULL,\n\n InvoiceDate DATETIME NOT NULL,\n\n BillingAddress NVARCHAR(70),\n\n BillingCity NVARCHAR(40),\n\n BillingState NVARCHAR(40),\n\n BillingCountry NVARCHAR(40),\n\n BillingPostalCode NVARCHAR(10),\n\n Total NUMERIC(10,2) NOT NULL,\n\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE INDEX IFK\_InvoiceCustomerId ON "invoices" (CustomerId)\n\nCREATE INDEX IFK\_InvoiceLineInvoiceId ON "invoice\_items" (InvoiceId)\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n InvoiceId INTEGER NOT NULL,\n\n TrackId INTEGER NOT NULL,\n\n UnitPrice NUMERIC(10,2) NOT NULL,\n\n Quantity INTEGER NOT NULL,\n\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE INDEX IFK\_InvoiceLineTrackId ON "invoice\_items" (TrackId)\n\nCREATE TABLE "customers"\n\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n FirstName NVARCHAR(40) NOT NULL,\n\n LastName NVARCHAR(20) NOT NULL,\n\n Company NVARCHAR(80),\n\n Address NVARCHAR(70),\n\n City NVARCHAR(40),\n\n State NVARCHAR(40),\n\n Country NVARCHAR(40),\n\n PostalCode NVARCHAR(10),\n\n Phone NVARCHAR(24),\n\n Fax NVARCHAR(24),\n\n Email NVARCHAR(60) NOT NULL,\n\n SupportRepId INTEGER,\n\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE INDEX IFK\_CustomerSupportRepId ON "customers" (SupportRepId)\n\n\n===Additional Context\n\n===Response Guidelines\n\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n\n3. If the provided context is insufficient, please explain why it can't be generated.\n\n4. Please use the most relevant table(s).\n\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n\n'}], {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': '\n\n List all albums and their corresponding artist names\n\n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName\n\nFROM albums\n\nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': '\n\n Find all tracks with a name containing "What" (case-insensitive)\n\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM tracks WHERE Name LIKE '%What%';'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': 'SELECT name FROM sqlite\_master WHERE type='table';'}, {'role': 'user', 'content': '\n\n Get the total number of invoices for each customer\n\n'}]


Using model gpt-4o-mini for 806.25 tokens (approx)

LLM Response: SELECT CustomerId, COUNT(\*) AS TotalInvoices  
FROM invoices  
GROUP BY CustomerId;  
Extracted SQL: SELECT CustomerId, COUNT(\*) AS TotalInvoices  
FROM invoices  
GROUP BY CustomerId;  
SELECT CustomerId, COUNT(\*) AS TotalInvoices

```
FROM invoices
GROUP BY CustomerId;
 CustomerId TotalInvoices
0 1 7
1 2 7
2 3 7
3 4 7
4 5 7
5 6 7
6 7 7
7 8 7
8 9 7
9 10 7
10 11 7
11 12 7
12 13 7
13 14 7
14 15 7
15 16 7
16 17 7
17 18 7
18 19 7
19 20 7
20 21 7
21 22 7
22 23 7
23 24 7
24 25 7
25 26 7
26 27 7
27 28 7
28 29 7
29 30 7
30 31 7
31 32 7
32 33 7
33 34 7
34 35 7
35 36 7
36 37 7
37 38 7
```

|    |    |   |
|----|----|---|
| 38 | 39 | 7 |
| 39 | 40 | 7 |
| 40 | 41 | 7 |
| 41 | 42 | 7 |
| 42 | 43 | 7 |
| 43 | 44 | 7 |
| 44 | 45 | 7 |
| 45 | 46 | 7 |
| 46 | 47 | 7 |
| 47 | 48 | 7 |
| 48 | 49 | 7 |
| 49 | 50 | 7 |
| 50 | 51 | 7 |
| 51 | 52 | 7 |
| 52 | 53 | 7 |
| 53 | 54 | 7 |
| 54 | 55 | 7 |
| 55 | 56 | 7 |
| 56 | 57 | 7 |
| 57 | 58 | 7 |
| 58 | 59 | 6 |

Using model gpt-4o-mini for 186.25 tokens (approx)



```
Out[25]: ('SELECT CustomerId, COUNT(*) AS TotalInvoices \nFROM invoices \nGROUP BY CustomerId;',
 CustomerId TotalInvoices
0 1 7
1 2 7
2 3 7
3 4 7
4 5 7
5 6 7
6 7 7
7 8 7
8 9 7
9 10 7
10 11 7
11 12 7
12 13 7
13 14 7
14 15 7
15 16 7
16 17 7
17 18 7
18 19 7
19 20 7
20 21 7
21 22 7
22 23 7
23 24 7
24 25 7
25 26 7
26 27 7
27 28 7
28 29 7
29 30 7
30 31 7
31 32 7
32 33 7
33 34 7
34 35 7
35 36 7
36 37 7
37 38 7
```



|    |    |    |
|----|----|----|
| 38 | 39 | 7  |
| 39 | 40 | 7  |
| 40 | 41 | 7  |
| 41 | 42 | 7  |
| 42 | 43 | 7  |
| 43 | 44 | 7  |
| 44 | 45 | 7  |
| 45 | 46 | 7  |
| 46 | 47 | 7  |
| 47 | 48 | 7  |
| 48 | 49 | 7  |
| 49 | 50 | 7  |
| 50 | 51 | 7  |
| 51 | 52 | 7  |
| 52 | 53 | 7  |
| 53 | 54 | 7  |
| 54 | 55 | 7  |
| 55 | 56 | 7  |
| 56 | 57 | 7  |
| 57 | 58 | 7  |
| 58 | 59 | 6, |

```
Figure({
 'data': [{
 'alignmentgroup': 'True',
 'hovernplate': 'Customer ID=%{x}
Total Invoices=%{y}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': '#636efa', 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,
 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36,
 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54,
 55, 56, 57, 58, 59]),
 'xaxis': 'x',
 'y': array([7, 7,
 7,
 7, 7, 7, 7, 7, 7, 7, 7, 7, 7, 6]),
 'yaxis': 'y'}],
 'xaxis': 'x',
 'yaxis': 'y'
})
```

```
'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Total Invoices per Customer'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'Customer ID'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Total Invoices'}}}
)))
```

```
In [26]: question = """
 Find the total number of invoices per country:
 """

 vn.ask(question=question)
```

```
Number of requested results 10 is greater than number of elements in index 6, updating n_results = 6
Number of requested results 10 is greater than number of elements in index 1, updating n_results = 1
```

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n(\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n CustomerId INTEGER NOT NULL,\n InvoiceDate DATETIME NOT NULL,\n BillingAddress NVARCHAR(70),\n BillingCity NVARCHAR(40),\n BillingState NVARCHAR(40),\n BillingCountry NVARCHAR(40),\n BillingPostalCode NVARCHAR(10),\n Total NUMERIC(10,2) NOT NULL,\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId)\n)\nON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE TABLE "invoice\_items"\n(\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n InvoiceId INTEGER NOT NULL,\n TrackId INTEGER NOT NULL,\n UnitPrice NUMERIC(10,2) NOT NULL,\n Quantity INTEGER NOT NULL,\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId)\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId)\n)\nON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE INDEX IFK\_InvoiceCustomerId ON "invoices" (CustomerId)\n\nCREATE INDEX IFK\_InvoiceLineInvoiceId ON "invoice\_items" (InvoiceId)\n\nCREATE INDEX IFK\_InvoiceLineTrackId ON "invoice\_items" (TrackId)\n\nCREATE TABLE "employees"\n(\n EmployeeId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n LastName NVARCHAR(20) NOT NULL,\n FirstName NVARCHAR(20) NOT NULL,\n Title NVARCHAR(30),\n ReportsTo INTEGER,\n BirthDate DATETIME,\n HireDate DATETIME,\n Address NVARCHAR(70),\n City NVARCHAR(40),\n State NVARCHAR(40),\n Country NVARCHAR(40),\n PostalCode NVARCHAR(10),\n Phone NVARCHAR(24),\n Fax NVARCHAR(24),\n Email NVARCHAR(60),\n FOREIGN KEY (ReportsTo) REFERENCES "employees" (EmployeeId)\n)\nON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE INDEX IFK\_EmployeeReportsTo ON "employees" (ReportsTo)\n\n\n===Additional Context\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'}], {'role': 'user', 'content': '\n Get the total number of invoices for each customer\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(\*) AS TotalInvoices\nFROM invoices\nGROUP BY CustomerId;'}, {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': '\n List all albums and their corresponding artist names\n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName\nFROM albums\nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': '\n Find all tracks with a name containing "What" (case-insensitive)\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM tracks WHERE Name LIKE '%What%';'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': 'SELECT name FROM sqlite\_master WHERE type='table';'}, {'role': 'user', 'content': '\n Find the total number of invoices per country:\n'}]

Using model gpt-4o-mini for 846.25 tokens (approx)

LLM Response: SELECT BillingCountry, COUNT(\*) AS TotalInvoices

FROM invoices

GROUP BY BillingCountry;


Extracted SQL: SELECT BillingCountry, COUNT(\*) AS TotalInvoices

FROM invoices

```
GROUP BY BillingCountry;
SELECT BillingCountry, COUNT(*) AS TotalInvoices
FROM invoices
GROUP BY BillingCountry;
```

|    | BillingCountry | TotalInvoices |
|----|----------------|---------------|
| 0  | Argentina      | 7             |
| 1  | Australia      | 7             |
| 2  | Austria        | 7             |
| 3  | Belgium        | 7             |
| 4  | Brazil         | 35            |
| 5  | Canada         | 56            |
| 6  | Chile          | 7             |
| 7  | Czech Republic | 14            |
| 8  | Denmark        | 7             |
| 9  | Finland        | 7             |
| 10 | France         | 35            |
| 11 | Germany        | 28            |
| 12 | Hungary        | 7             |
| 13 | India          | 13            |
| 14 | Ireland        | 7             |
| 15 | Italy          | 7             |
| 16 | Netherlands    | 7             |
| 17 | Norway         | 7             |
| 18 | Poland         | 7             |
| 19 | Portugal       | 14            |
| 20 | Spain          | 7             |
| 21 | Sweden         | 7             |
| 22 | USA            | 91            |
| 23 | United Kingdom | 21            |

Using model gpt-4o-mini for 188.25 tokens (approx)



Out[26]: ('SELECT BillingCountry, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY BillingCountry;',

|    | BillingCountry | TotalInvoices |
|----|----------------|---------------|
| 0  | Argentina      | 7             |
| 1  | Australia      | 7             |
| 2  | Austria        | 7             |
| 3  | Belgium        | 7             |
| 4  | Brazil         | 35            |
| 5  | Canada         | 56            |
| 6  | Chile          | 7             |
| 7  | Czech Republic | 14            |
| 8  | Denmark        | 7             |
| 9  | Finland        | 7             |
| 10 | France         | 35            |
| 11 | Germany        | 28            |
| 12 | Hungary        | 7             |
| 13 | India          | 13            |
| 14 | Ireland        | 7             |
| 15 | Italy          | 7             |
| 16 | Netherlands    | 7             |
| 17 | Norway         | 7             |
| 18 | Poland         | 7             |
| 19 | Portugal       | 14            |
| 20 | Spain          | 7             |
| 21 | Sweden         | 7             |
| 22 | USA            | 91            |
| 23 | United Kingdom | 21,           |

```
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovernplate': 'Country=%{x}
Total Invoices=%{marker.color}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': array([7, 7, 7, 7, 35, 56, 7, 14, 7, 7, 35, 28, 7, 13, 7, 7, 7, 7,
 7, 14, 7, 7, 91, 21]),
 'coloraxis': 'coloraxis',
 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
```

```

 'x': array(['Argentina', 'Australia', 'Austria', 'Belgium', 'Brazil', 'Canada',
 'Chile', 'Czech Republic', 'Denmark', 'Finland', 'France', 'Germany',
 'Hungary', 'India', 'Ireland', 'Italy', 'Netherlands', 'Norway',
 'Poland', 'Portugal', 'Spain', 'Sweden', 'USA', 'United Kingdom'],
 dtype=object),
 'xaxis': 'x',
 'y': array([7, 7, 7, 7, 35, 56, 7, 14, 7, 7, 35, 28, 7, 13, 7, 7, 7, 7,
 7, 14, 7, 7, 91, 21]),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'coloraxis': {'colorbar': {'title': {'text': 'Total Invoices'}}},
 'colorscale': [[0.0, '#440154'], [0.11111111111111111,
 '#482878'], [0.22222222222222222,
 '#3e4989'], [0.33333333333333333,
 '#31688e'], [0.44444444444444444,
 '#26828e'], [0.55555555555555556,
 '#1f9e89'], [0.66666666666666666,
 '#35b779'], [0.77777777777777778,
 '#6ece58'], [0.88888888888888888,
 '#b5de2b'], [1.0, '#fde725']]}],
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Total Number of Invoices per Country'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'Country'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Total Invoices'}}}
)))

```

```

In [27]: question = """
 List all invoices with a total exceeding $10:
 """

 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 7, updating n\_results = 7  
 Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoice\_items"\n(\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n InvoiceId INTEGER NOT NULL,\n TrackId INTEGER NOT NULL,\n UnitPrice NUMERIC(10,2) NOT NULL,\n Quantity INTEGER NOT NULL,\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId)\n)\nDELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId)\n)\nDELETE NO ACTION ON UPDATE NO ACTION\n)\nCREATE INDEX IFK\_InvoiceLineInvoiceId ON "invoice\_items" (InvoiceId)\nCREATE TABLE "invoices"\n(\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n CustomerId INTEGER NOT NULL,\n InvoiceDate DATETIME NOT NULL,\n BillingAddress NVARCHAR(70),\n BillingCity NVARCHAR(40),\n BillingState NVARCHAR(40),\n BillingCountry NVARCHAR(40),\n BillingPostalCode NVARCHAR(10),\n Total NUMERIC(10,2) NOT NULL,\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId)\n)\nDELETE NO ACTION ON UPDATE NO ACTION\n)\nCREATE INDEX IFK\_InvoiceLineTrackId ON "invoice\_items" (TrackId)\nCREATE INDEX IFK\_InvoiceCustomerId ON "invoices" (CustomerId)\nCREATE TABLE "tracks"\n(\n TrackId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n Name NVARCHAR(200) NOT NULL,\n AlbumId INTEGER,\n MediaTypeId INTEGER NOT NULL,\n GenreId INTEGER,\n Composer NVARCHAR(220),\n Milliseconds INTEGER NOT NULL,\n Bytes INTEGER,\n UnitPrice NUMERIC(10,2) NOT NULL,\n FOREIGN KEY (AlbumId) REFERENCES "albums" (AlbumId)\n)\nDELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (GenreId) REFERENCES "genres" (GenreId)\n)\nDELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (MediaTypeId) REFERENCES "media\_types" (MediaTypeId)\n)\nDELETE NO ACTION ON UPDATE NO ACTION\n)\n\n===Additional Context\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'}], [{'role': 'user', 'content': '\n Get the total number of invoices for each customer\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(\*) AS TotalInvoices\nFROM invoices\nGROUP BY CustomerId;'}, {'role': 'user', 'content': '\n Find the total number of invoices per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(\*) AS TotalInvoices\nFROM invoices\nGROUP BY BillingCountry;'}, {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': '\n List all albums and their corresponding artist names\n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName\nFROM albums\nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': '\n Find all tracks with a name containing "What" (case-insensitive)\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM tracks WHERE Name LIKE '%What%';'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': 'SELECT name FROM sqlite\_master WHERE type='table';'}, {'role': 'user', 'content': '\n List all invoices with a total exceeding \$10:\n'}]

Using model gpt-4o-mini for 881.75 tokens (approx)

LLM Response: SELECT \* FROM invoices WHERE Total > 10;

Extracted SQL: SELECT \* FROM invoices WHERE Total > 10;

SELECT \* FROM invoices WHERE Total > 10;



|    | InvoiceId | CustomerId | InvoiceDate         | BillingAddress \          |
|----|-----------|------------|---------------------|---------------------------|
| 0  | 5         | 23         | 2009-01-11 00:00:00 | 69 Salem Street           |
| 1  | 12        | 2          | 2009-02-11 00:00:00 | Theodor-Heuss-Straße 34   |
| 2  | 19        | 40         | 2009-03-14 00:00:00 | 8, Rue Hanovre            |
| 3  | 26        | 19         | 2009-04-14 00:00:00 | 1 Infinite Loop           |
| 4  | 33        | 57         | 2009-05-15 00:00:00 | Calle Lira, 198           |
| .. | ...       | ...        | ...                 | ...                       |
| 59 | 383       | 10         | 2013-08-12 00:00:00 | Rua Dr. Falcão Filho, 155 |
| 60 | 390       | 48         | 2013-09-12 00:00:00 | Lijnbaansgracht 120bg     |
| 61 | 397       | 27         | 2013-10-13 00:00:00 | 1033 N Park Ave           |
| 62 | 404       | 6          | 2013-11-13 00:00:00 | Rilská 3174/6             |
| 63 | 411       | 44         | 2013-12-14 00:00:00 | Porthaninkatu 9           |

|    | BillingCity | BillingState | BillingCountry | BillingPostalCode | Total |
|----|-------------|--------------|----------------|-------------------|-------|
| 0  | Boston      | MA           | USA            | 2113              | 13.86 |
| 1  | Stuttgart   | None         | Germany        | 70174             | 13.86 |
| 2  | Paris       | None         | France         | 75002             | 13.86 |
| 3  | Cupertino   | CA           | USA            | 95014             | 13.86 |
| 4  | Santiago    | None         | Chile          | None              | 13.86 |
| .. | ...         | ...          | ...            | ...               | ...   |
| 59 | São Paulo   | SP           | Brazil         | 01007-010         | 13.86 |
| 60 | Amsterdam   | VV           | Netherlands    | 1016              | 13.86 |
| 61 | Tucson      | AZ           | USA            | 85719             | 13.86 |
| 62 | Prague      | None         | Czech Republic | 14300             | 25.86 |
| 63 | Helsinki    | None         | Finland        | 00530             | 13.86 |

[64 rows x 9 columns]

Using model gpt-4o-mini for 228.5 tokens (approx)

# Invoices with Total Exceeding \$10



```
Out[27]: ('SELECT * FROM invoices WHERE Total > 10;',
```

|    | InvoiceId | CustomerId | InvoiceDate         | BillingAddress \          |
|----|-----------|------------|---------------------|---------------------------|
| 0  | 5         | 23         | 2009-01-11 00:00:00 | 69 Salem Street           |
| 1  | 12        | 2          | 2009-02-11 00:00:00 | Theodor-Heuss-Straße 34   |
| 2  | 19        | 40         | 2009-03-14 00:00:00 | 8, Rue Hanovre            |
| 3  | 26        | 19         | 2009-04-14 00:00:00 | 1 Infinite Loop           |
| 4  | 33        | 57         | 2009-05-15 00:00:00 | Calle Lira, 198           |
| .. | ...       | ...        | ...                 | ...                       |
| 59 | 383       | 10         | 2013-08-12 00:00:00 | Rua Dr. Falcão Filho, 155 |
| 60 | 390       | 48         | 2013-09-12 00:00:00 | Lijnbaansgracht 120bg     |
| 61 | 397       | 27         | 2013-10-13 00:00:00 | 1033 N Park Ave           |
| 62 | 404       | 6          | 2013-11-13 00:00:00 | Rilská 3174/6             |
| 63 | 411       | 44         | 2013-12-14 00:00:00 | Porthaninkatu 9           |

|    | BillingCity | BillingState | BillingCountry | BillingPostalCode | Total |
|----|-------------|--------------|----------------|-------------------|-------|
| 0  | Boston      | MA           | USA            | 2113              | 13.86 |
| 1  | Stuttgart   | None         | Germany        | 70174             | 13.86 |
| 2  | Paris       | None         | France         | 75002             | 13.86 |
| 3  | Cupertino   | CA           | USA            | 95014             | 13.86 |
| 4  | Santiago    | None         | Chile          | None              | 13.86 |
| .. | ...         | ...          | ...            | ...               | ...   |
| 59 | São Paulo   | SP           | Brazil         | 01007-010         | 13.86 |
| 60 | Amsterdam   | VV           | Netherlands    | 1016              | 13.86 |
| 61 | Tucson      | AZ           | USA            | 85719             | 13.86 |
| 62 | Prague      | None         | Czech Republic | 14300             | 25.86 |
| 63 | Helsinki    | None         | Finland        | 00530             | 13.86 |

```
[64 rows x 9 columns],
```

```
Figure({
```

```
 'data': [{'alignmentgroup': 'True',
 'hovernplate': 'InvoiceId=%{x}
Total=%{y}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': '#636efa', 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array([5, 12, 19, 26, 33, 40, 47, 54, 61, 68, 75, 82, 88, 89,
```

```

 96, 103, 110, 117, 124, 131, 138, 145, 152, 159, 166, 173, 180, 187,
 193, 194, 201, 208, 215, 222, 229, 236, 243, 250, 257, 264, 271, 278,
 285, 292, 298, 299, 306, 311, 312, 313, 320, 327, 334, 341, 348, 355,
 362, 369, 376, 383, 390, 397, 404, 411]),
 'xaxis': 'x',
 'y': array([13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86,
 13.86, 13.86, 17.91, 18.86, 21.86, 15.86, 13.86, 13.86, 13.86, 13.86,
 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 14.91, 21.86,
 18.86, 15.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86,
 13.86, 13.86, 13.86, 13.86, 10.91, 23.86, 16.86, 11.94, 10.91, 16.86,
 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86, 13.86,
 13.86, 13.86, 25.86, 13.86]),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Invoices with Total Exceeding $10'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'InvoiceId'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Total'}}}
)))

```

```

In [28]: question = """
 Find all invoices since 2010 and the total amount invoiced:
 """

 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 8, updating n\_results = 8  
 Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

[Downloads/openai-gpt-4o-mini-chromadb-sqlite-test-1.html](#)

WHERE InvoiceDate >= '2010-01-01';

Extracted SQL: SELECT \*, Total

FROM invoices

WHERE InvoiceDate >= '2010-01-01';

SELECT \*, Total

FROM invoices

WHERE InvoiceDate >= '2010-01-01';

|     | InvoiceId | CustomerId | InvoiceDate         | \ |
|-----|-----------|------------|---------------------|---|
| 0   | 84        | 43         | 2010-01-08 00:00:00 |   |
| 1   | 85        | 45         | 2010-01-08 00:00:00 |   |
| 2   | 86        | 47         | 2010-01-09 00:00:00 |   |
| 3   | 87        | 51         | 2010-01-10 00:00:00 |   |
| 4   | 88        | 57         | 2010-01-13 00:00:00 |   |
| ..  | ...       | ...        | ...                 |   |
| 324 | 408       | 25         | 2013-12-05 00:00:00 |   |
| 325 | 409       | 29         | 2013-12-06 00:00:00 |   |
| 326 | 410       | 35         | 2013-12-09 00:00:00 |   |
| 327 | 411       | 44         | 2013-12-14 00:00:00 |   |
| 328 | 412       | 58         | 2013-12-22 00:00:00 |   |

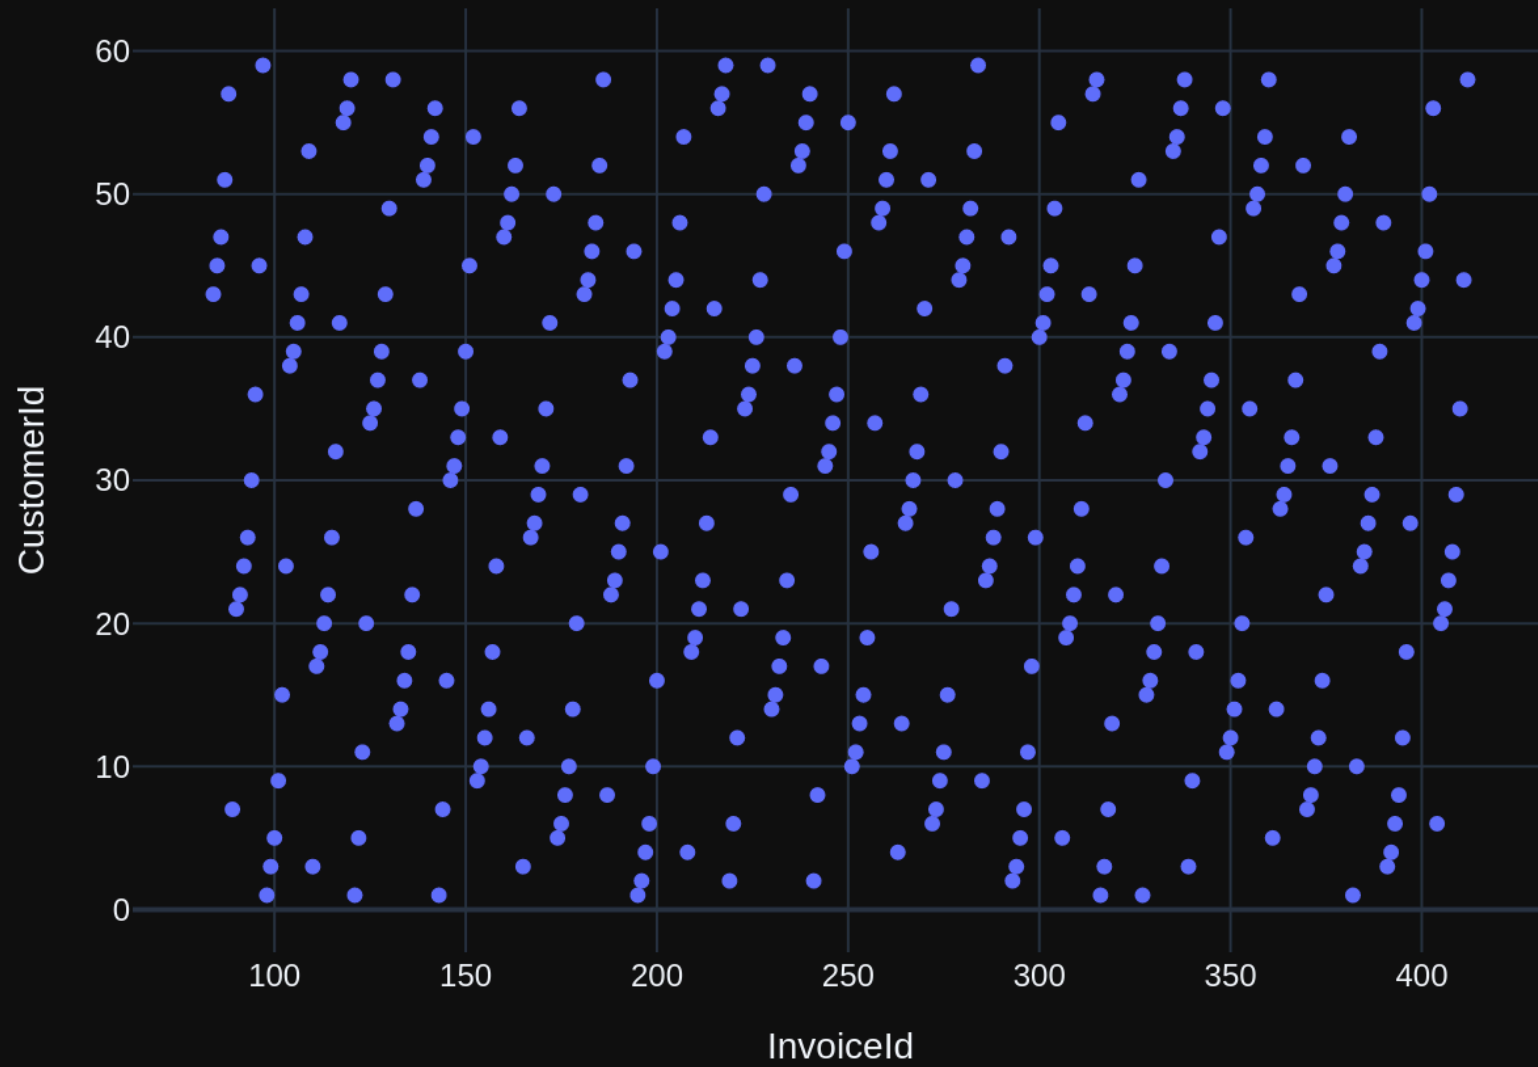
|     | BillingAddress                           | BillingCity | BillingState | \ |
|-----|------------------------------------------|-------------|--------------|---|
| 0   | 68, Rue Jouvence                         | Dijon       | None         |   |
| 1   | Erzsébet krt. 58.                        | Budapest    | None         |   |
| 2   | Via Degli Scipioni, 43                   | Rome        | RM           |   |
| 3   | Celsiusg. 9                              | Stockholm   | None         |   |
| 4   | Calle Lira, 198                          | Santiago    | None         |   |
| ..  | ...                                      | ...         | ...          |   |
| 324 | 319 N. Frances Street                    | Madison     | WI           |   |
| 325 | 796 Dundas Street West                   | Toronto     | ON           |   |
| 326 | Rua dos Campeões Europeus de Viena, 4350 | Porto       | None         |   |
| 327 | Porthaninkatu 9                          | Helsinki    | None         |   |
| 328 | 12,Community Centre                      | Delhi       | None         |   |

|     | BillingCountry | BillingPostalCode | Total | Total |
|-----|----------------|-------------------|-------|-------|
| 0   | France         | 21000             | 1.98  | 1.98  |
| 1   | Hungary        | H-1073            | 1.98  | 1.98  |
| 2   | Italy          | 00192             | 3.96  | 3.96  |
| 3   | Sweden         | 11230             | 6.94  | 6.94  |
| 4   | Chile          | None              | 17.91 | 17.91 |
| ..  | ...            | ...               | ...   | ...   |
| 324 | USA            | 53703             | 3.96  | 3.96  |

|     |          |         |       |       |
|-----|----------|---------|-------|-------|
| 325 | Canada   | M6J 1V1 | 5.94  | 5.94  |
| 326 | Portugal | None    | 8.91  | 8.91  |
| 327 | Finland  | 00530   | 13.86 | 13.86 |
| 328 | India    | 110017  | 1.99  | 1.99  |

[329 rows x 10 columns]

Using model gpt-4o-mini for 245.75 tokens (approx)





Out[28]: ("SELECT \*, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';",

|     | InvoiceId | CustomerId | InvoiceDate \       |
|-----|-----------|------------|---------------------|
| 0   | 84        | 43         | 2010-01-08 00:00:00 |
| 1   | 85        | 45         | 2010-01-08 00:00:00 |
| 2   | 86        | 47         | 2010-01-09 00:00:00 |
| 3   | 87        | 51         | 2010-01-10 00:00:00 |
| 4   | 88        | 57         | 2010-01-13 00:00:00 |
| ..  | ...       | ...        | ...                 |
| 324 | 408       | 25         | 2013-12-05 00:00:00 |
| 325 | 409       | 29         | 2013-12-06 00:00:00 |
| 326 | 410       | 35         | 2013-12-09 00:00:00 |
| 327 | 411       | 44         | 2013-12-14 00:00:00 |
| 328 | 412       | 58         | 2013-12-22 00:00:00 |

|     | BillingAddress                           | BillingCity | BillingState \ |
|-----|------------------------------------------|-------------|----------------|
| 0   | 68, Rue Jouvence                         | Dijon       | None           |
| 1   | Erzsébet krt. 58.                        | Budapest    | None           |
| 2   | Via Degli Scipioni, 43                   | Rome        | RM             |
| 3   | Celsiusg. 9                              | Stockholm   | None           |
| 4   | Calle Lira, 198                          | Santiago    | None           |
| ..  | ...                                      | ...         | ...            |
| 324 | 319 N. Frances Street                    | Madison     | WI             |
| 325 | 796 Dundas Street West                   | Toronto     | ON             |
| 326 | Rua dos Campeões Europeus de Viena, 4350 | Porto       | None           |
| 327 | Porthaninkatu 9                          | Helsinki    | None           |
| 328 | 12,Community Centre                      | Delhi       | None           |

|     | BillingCountry | BillingPostalCode | Total | Total |
|-----|----------------|-------------------|-------|-------|
| 0   | France         | 21000             | 1.98  | 1.98  |
| 1   | Hungary        | H-1073            | 1.98  | 1.98  |
| 2   | Italy          | 00192             | 3.96  | 3.96  |
| 3   | Sweden         | 11230             | 6.94  | 6.94  |
| 4   | Chile          | None              | 17.91 | 17.91 |
| ..  | ...            | ...               | ...   | ...   |
| 324 | USA            | 53703             | 3.96  | 3.96  |
| 325 | Canada         | M6J 1V1           | 5.94  | 5.94  |
| 326 | Portugal       | None              | 8.91  | 8.91  |
| 327 | Finland        | 00530             | 13.86 | 13.86 |
| 328 | India          | 110017            | 1.99  | 1.99  |

```
[329 rows x 10 columns],
Figure({
 'data': [{ 'hovertemplate': 'InvoiceId=%{x}
CustomerId=%{y}<extra></extra>',
 'legendgroup': '',
 'marker': { 'color': '#636efa', 'symbol': 'circle' },
 'mode': 'markers',
 'name': '',
 'orientation': 'v',
 'showlegend': False,
 'type': 'scatter',
 'x': array([84, 85, 86, ..., 410, 411, 412]),
 'xaxis': 'x',
 'y': array([43, 45, 47, ..., 35, 44, 58]),
 'yaxis': 'y' }],
 'layout': { 'legend': { 'tracegroupgap': 0 },
 'margin': { 't': 60 },
 'template': '...',
 'xaxis': { 'anchor': 'y', 'domain': [0.0, 1.0], 'title': { 'text': 'InvoiceId' } },
 'yaxis': { 'anchor': 'x', 'domain': [0.0, 1.0], 'title': { 'text': 'CustomerId' } } }
})
```

```
In [29]: question = """
 List all employees and their reporting manager's name (if any):
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 9, updating n\_results = 9  
 Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables \nCREATE INDEX IFK\_EmployeeReportsTo ON "employees" (ReportsTo)\n\nCREATE TABLE "employees" \n\n(\n EmployeeId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n LastName NVARCHAR(20) NOT NULL,\n FirstName NVARCHAR(20) NOT NULL,\n Title NVARCHAR(30),\n ReportsTo INTEGER,\n BirthDate DATETIME,\n HireDate DATETIME,\n Address NVARCHAR(70),\n City NVARCHAR(40),\n State NVARCHAR(40),\n Country NVARCHAR(40),\n PostalCode NVARCHAR(10),\n Phone NVARCHAR(24),\n Fax NVARCHAR(24),\n Email NVARCHAR(60),\n FOREIGN KEY (ReportsTo) REFERENCES "employees" (EmployeeId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n)\n\nCREATE TABLE "customers" \n\n(\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n FirstName NVARCHAR(40) NOT NULL,\n LastName NVARCHAR(20) NOT NULL,\n Company NVARCHAR(80),\n Address NVARCHAR(70),\n City NVARCHAR(40),\n State NVARCHAR(40),\n Country NVARCHAR(40),\n PostalCode NVARCHAR(10),\n Phone NVARCHAR(24),\n Fax NVARCHAR(24),\n Email NVARCHAR(60) NOT NULL,\n SupportRepId INTEGER,\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n)\n\nCREATE INDEX IFK\_CustomerSupportRepId ON "customers" (SupportRepId)\n\nCREATE TABLE "invoices" \n\n(\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n CustomerId INTEGER NOT NULL,\n InvoiceDate DATETIME NOT NULL,\n BillingAddress NVARCHAR(70),\n BillingCity NVARCHAR(40),\n BillingState NVARCHAR(40),\n BillingCountry NVARCHAR(40),\n BillingPostalCode NVARCHAR(10),\n Total NUMERIC(10, 2) NOT NULL,\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n)\n\nCREATE TABLE sqlite\_stat1(tbl,idx,stat)\n\n\n===Additional Context \n\n===Response Guidelines \n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question. \n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql \n3. If the provided context is insufficient, please explain why it can't be generated. \n4. Please use the most relevant table(s). \n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before. \n'}], {'role': 'user', 'content': ' \n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': "SELECT \*, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n Get the total number of invoices for each customer\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY CustomerId;'}, {'role': 'user', 'content': ' \n List all albums and their corresponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': ' \n Find the total number of invoices per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY BillingCountry;'}, {'role': 'user', 'content': ' \n List all invoices with a total exceeding \$10:\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': ' \n Find all tracks with a name containing "What" (case-insensitive)\n'}, {'role': 'assistant', 'content': "SELECT \* FROM tracks WHERE Name LIKE '%What%'"}, {'role': 'user', 'content': ' \n List all employees and their reporting manager's name (if any):\n"}]

Using model gpt-4o-mini for 946.5 tokens (approx)

LLM Response: SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName,  
m.FirstName AS ManagerFirstName, m.LastName AS ManagerLastName

FROM employees e

LEFT JOIN employees m ON e.ReportsTo = m.EmployeeId;

Extracted SQL: SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName,  
m.FirstName AS ManagerFirstName, m.LastName AS ManagerLastName

FROM employees e

LEFT JOIN employees m ON e.ReportsTo = m.EmployeeId;

SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName,  
m.FirstName AS ManagerFirstName, m.LastName AS ManagerLastName

FROM employees e

LEFT JOIN employees m ON e.ReportsTo = m.EmployeeId;

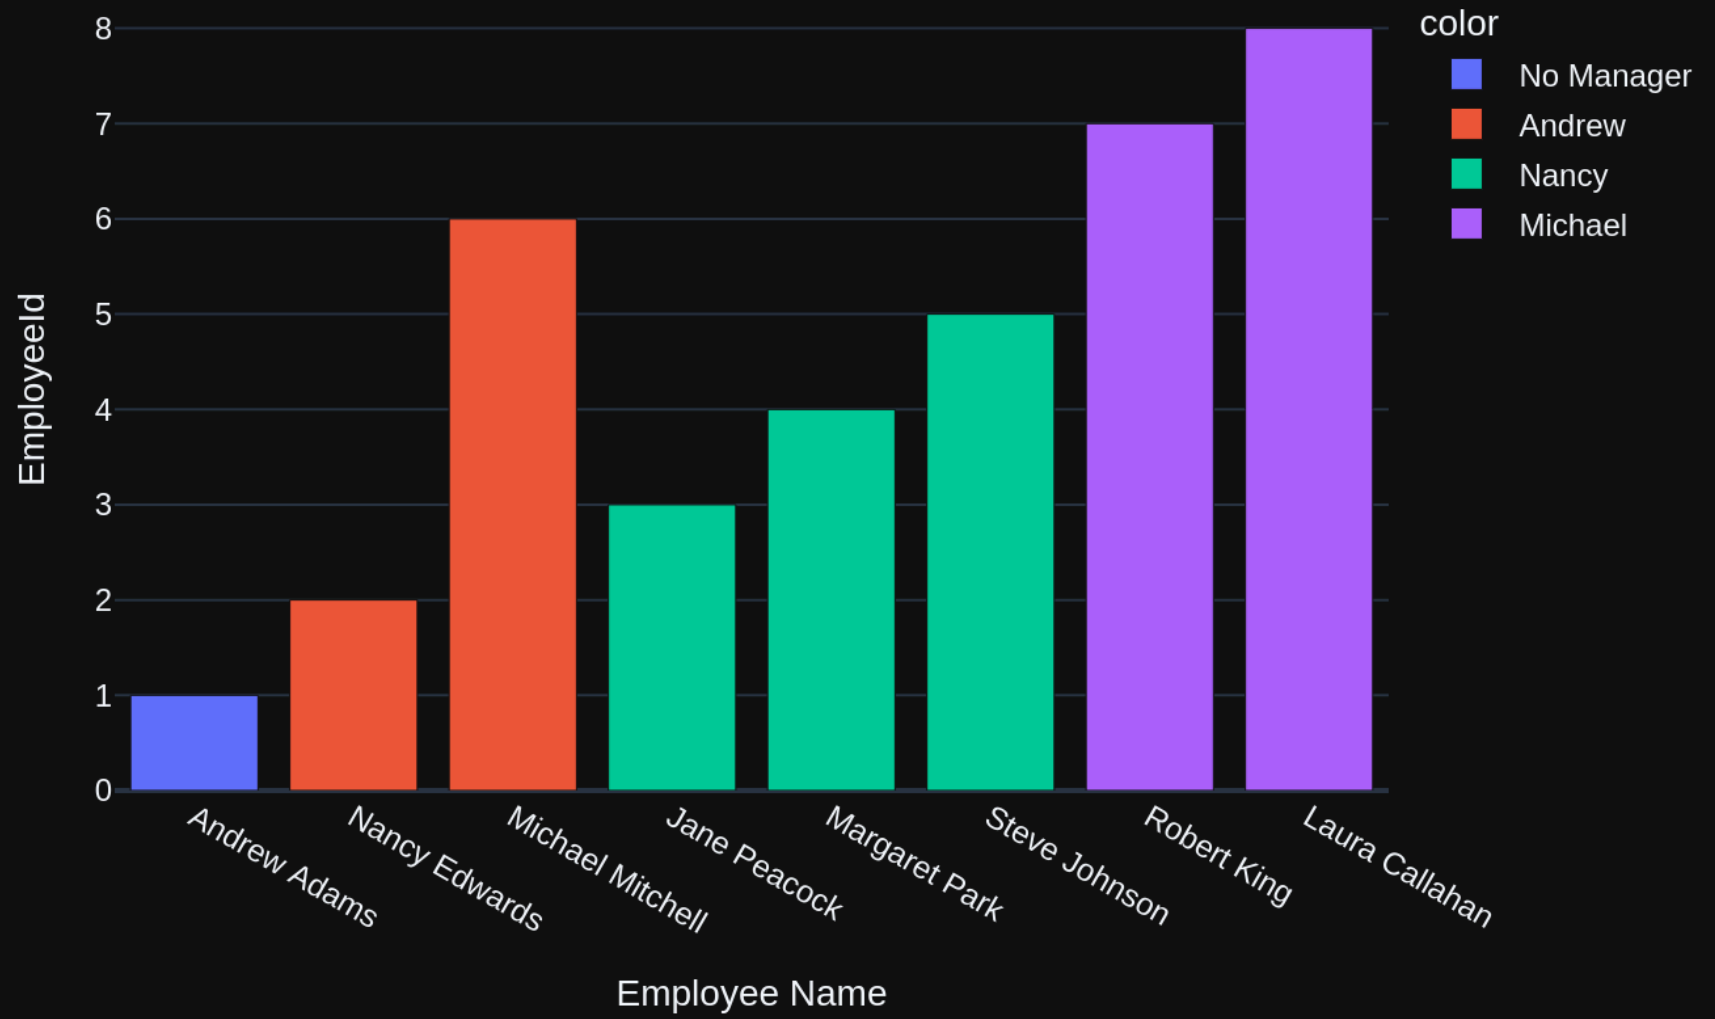
|   | EmployeeId | EmployeeFirstName | EmployeeLastName | ManagerFirstName | \ |
|---|------------|-------------------|------------------|------------------|---|
| 0 | 1          | Andrew            | Adams            | None             |   |
| 1 | 2          | Nancy             | Edwards          | Andrew           |   |
| 2 | 3          | Jane              | Peacock          | Nancy            |   |
| 3 | 4          | Margaret          | Park             | Nancy            |   |
| 4 | 5          | Steve             | Johnson          | Nancy            |   |
| 5 | 6          | Michael           | Mitchell         | Andrew           |   |
| 6 | 7          | Robert            | King             | Michael          |   |
| 7 | 8          | Laura             | Callahan         | Michael          |   |

ManagerLastName

|   |          |
|---|----------|
| 0 | None     |
| 1 | Adams    |
| 2 | Edwards  |
| 3 | Edwards  |
| 4 | Edwards  |
| 5 | Adams    |
| 6 | Mitchell |
| 7 | Mitchell |

Using model gpt-4o-mini for 249.5 tokens (approx)

## Employees and Their Managers



```
Out[29]: ('SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName, \n m.FirstName AS ManagerFirstName, m.LastName AS ManagerLastName\nFROM employees e\nLEFT JOIN employees m ON e.ReportsTo = m.EmployeeId;',
```

|   | EmployeeId | EmployeeFirstName | EmployeeLastName | ManagerFirstName | \ |
|---|------------|-------------------|------------------|------------------|---|
| 0 | 1          | Andrew            | Adams            | None             |   |
| 1 | 2          | Nancy             | Edwards          | Andrew           |   |
| 2 | 3          | Jane              | Peacock          | Nancy            |   |
| 3 | 4          | Margaret          | Park             | Nancy            |   |
| 4 | 5          | Steve             | Johnson          | Nancy            |   |
| 5 | 6          | Michael           | Mitchell         | Andrew           |   |
| 6 | 7          | Robert            | King             | Michael          |   |
| 7 | 8          | Laura             | Callahan         | Michael          |   |

|   | ManagerLastName |
|---|-----------------|
| 0 | None            |
| 1 | Adams           |
| 2 | Edwards         |
| 3 | Edwards         |
| 4 | Edwards         |
| 5 | Adams           |
| 6 | Mitchell        |
| 7 | Mitchell        |

```
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovernment': 'color=No Manager
Employee Name=%{x}
EmployeeId=%{y}<extra></extra>',
 'legendgroup': 'No Manager',
 'marker': {'color': '#636efa', 'pattern': {'shape': ''}},
 'name': 'No Manager',
 'offsetgroup': 'No Manager',
 'orientation': 'v',
 'showlegend': True,
 'textposition': 'auto',
 'type': 'bar',
 'x': array(['Andrew Adams'], dtype=object),
 'xaxis': 'x',
 'y': array([1]),
 'yaxis': 'y'},
 {'alignmentgroup': 'True',
 'hovernment': 'color=Andrew
Employee Name=%{x}
EmployeeId=%{y}<extra></extra>',
 'legendgroup': 'Andrew',
```

```

'marker': {'color': '#EF553B', 'pattern': {'shape': ''}},
'name': 'Andrew',
'offsetgroup': 'Andrew',
'orientation': 'v',
'showlegend': True,
'textposition': 'auto',
'type': 'bar',
'x': array(['Nancy Edwards', 'Michael Mitchell'], dtype=object),
'xaxis': 'x',
'y': array([2, 6]),
'yaxis': 'y'},
{'alignmentgroup': 'True',
'hovertemplate': 'color=Nancy
Employee Name=%{x}
EmployeeId=%{y}<extra></extra>',
'legendgroup': 'Nancy',
'marker': {'color': '#00cc96', 'pattern': {'shape': ''}},
'name': 'Nancy',
'offsetgroup': 'Nancy',
'orientation': 'v',
'showlegend': True,
'textposition': 'auto',
'type': 'bar',
'x': array(['Jane Peacock', 'Margaret Park', 'Steve Johnson'], dtype=object),
'xaxis': 'x',
'y': array([3, 4, 5]),
'yaxis': 'y'},
{'alignmentgroup': 'True',
'hovertemplate': 'color=Michael
Employee Name=%{x}
EmployeeId=%{y}<extra></extra>',
'legendgroup': 'Michael',
'marker': {'color': '#ab63fa', 'pattern': {'shape': ''}},
'name': 'Michael',
'offsetgroup': 'Michael',
'orientation': 'v',
'showlegend': True,
'textposition': 'auto',
'type': 'bar',
'x': array(['Robert King', 'Laura Callahan'], dtype=object),
'xaxis': 'x',
'y': array([7, 8]),
'yaxis': 'y'}],
'layout': {'barmode': 'relative',
'legend': {'title': {'text': 'color'}, 'tracegroupgap': 0},

```

```
 'template': '...',
 'title': {'text': 'Employees and Their Managers'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'Employee Name'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'EmployeeId'}}}
))
```

```
In [30]: question = """
 Get the average invoice total for each customer:
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1



65/129

ole': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': ' \n Get the average invoice total for each customer:\n'}]

Using model gpt-4o-mini for 1017.5 tokens (approx)

LLM Response: SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal

FROM invoices

GROUP BY CustomerId;

Extracted SQL: SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal

FROM invoices

GROUP BY CustomerId;

SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal

FROM invoices

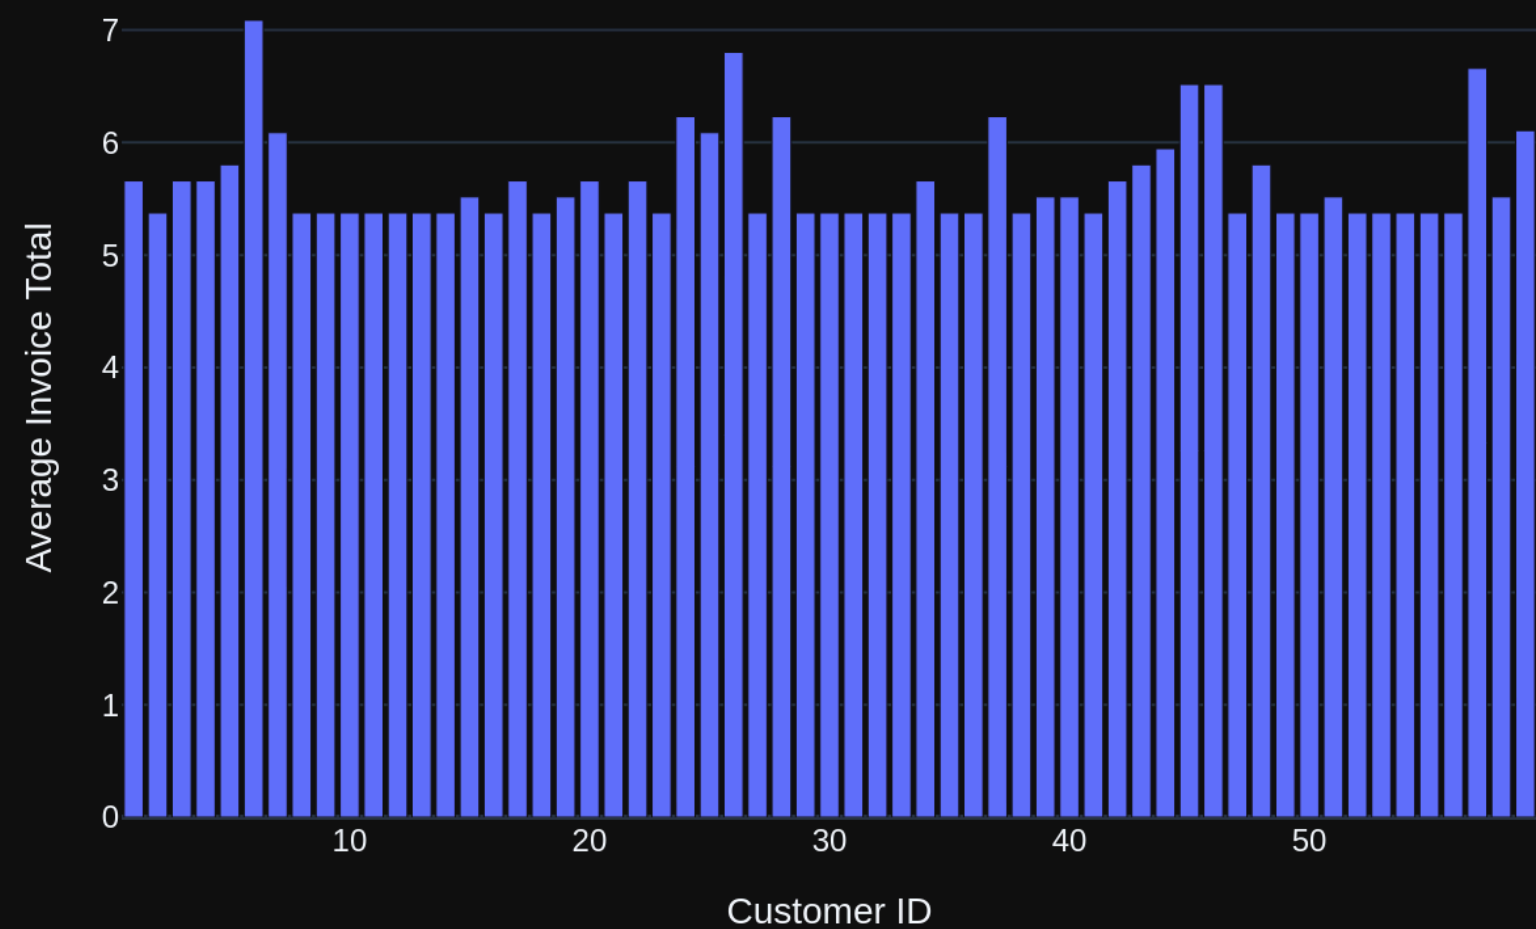
GROUP BY CustomerId;

|    | CustomerId | AverageInvoiceTotal |
|----|------------|---------------------|
| 0  | 1          | 5.660000            |
| 1  | 2          | 5.374286            |
| 2  | 3          | 5.660000            |
| 3  | 4          | 5.660000            |
| 4  | 5          | 5.802857            |
| 5  | 6          | 7.088571            |
| 6  | 7          | 6.088571            |
| 7  | 8          | 5.374286            |
| 8  | 9          | 5.374286            |
| 9  | 10         | 5.374286            |
| 10 | 11         | 5.374286            |
| 11 | 12         | 5.374286            |
| 12 | 13         | 5.374286            |
| 13 | 14         | 5.374286            |
| 14 | 15         | 5.517143            |
| 15 | 16         | 5.374286            |
| 16 | 17         | 5.660000            |
| 17 | 18         | 5.374286            |
| 18 | 19         | 5.517143            |
| 19 | 20         | 5.660000            |
| 20 | 21         | 5.374286            |
| 21 | 22         | 5.660000            |
| 22 | 23         | 5.374286            |
| 23 | 24         | 6.231429            |
| 24 | 25         | 6.088571            |
| 25 | 26         | 6.802857            |
| 26 | 27         | 5.374286            |

|    |    |          |
|----|----|----------|
| 27 | 28 | 6.231429 |
| 28 | 29 | 5.374286 |
| 29 | 30 | 5.374286 |
| 30 | 31 | 5.374286 |
| 31 | 32 | 5.374286 |
| 32 | 33 | 5.374286 |
| 33 | 34 | 5.660000 |
| 34 | 35 | 5.374286 |
| 35 | 36 | 5.374286 |
| 36 | 37 | 6.231429 |
| 37 | 38 | 5.374286 |
| 38 | 39 | 5.517143 |
| 39 | 40 | 5.517143 |
| 40 | 41 | 5.374286 |
| 41 | 42 | 5.660000 |
| 42 | 43 | 5.802857 |
| 43 | 44 | 5.945714 |
| 44 | 45 | 6.517143 |
| 45 | 46 | 6.517143 |
| 46 | 47 | 5.374286 |
| 47 | 48 | 5.802857 |
| 48 | 49 | 5.374286 |
| 49 | 50 | 5.374286 |
| 50 | 51 | 5.517143 |
| 51 | 52 | 5.374286 |
| 52 | 53 | 5.374286 |
| 53 | 54 | 5.374286 |
| 54 | 55 | 5.374286 |
| 55 | 56 | 5.374286 |
| 56 | 57 | 6.660000 |
| 57 | 58 | 5.517143 |
| 58 | 59 | 6.106667 |

Using model gpt-4o-mini for 191.75 tokens (approx)

Average Invoice Total for Each Customer



```
Out[30]: ('SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal \nFROM invoices \nGROUP BY CustomerId;',
 CustomerId AverageInvoiceTotal
0 1 5.660000
1 2 5.374286
2 3 5.660000
3 4 5.660000
4 5 5.802857
5 6 7.088571
6 7 6.088571
7 8 5.374286
8 9 5.374286
9 10 5.374286
10 11 5.374286
11 12 5.374286
12 13 5.374286
13 14 5.374286
14 15 5.517143
15 16 5.374286
16 17 5.660000
17 18 5.374286
18 19 5.517143
19 20 5.660000
20 21 5.374286
21 22 5.660000
22 23 5.374286
23 24 6.231429
24 25 6.088571
25 26 6.802857
26 27 5.374286
27 28 6.231429
28 29 5.374286
29 30 5.374286
30 31 5.374286
31 32 5.374286
32 33 5.374286
33 34 5.660000
34 35 5.374286
35 36 5.374286
36 37 6.231429
37 38 5.374286
```

|    |    |           |
|----|----|-----------|
| 38 | 39 | 5.517143  |
| 39 | 40 | 5.517143  |
| 40 | 41 | 5.374286  |
| 41 | 42 | 5.660000  |
| 42 | 43 | 5.802857  |
| 43 | 44 | 5.945714  |
| 44 | 45 | 6.517143  |
| 45 | 46 | 6.517143  |
| 46 | 47 | 5.374286  |
| 47 | 48 | 5.802857  |
| 48 | 49 | 5.374286  |
| 49 | 50 | 5.374286  |
| 50 | 51 | 5.517143  |
| 51 | 52 | 5.374286  |
| 52 | 53 | 5.374286  |
| 53 | 54 | 5.374286  |
| 54 | 55 | 5.374286  |
| 55 | 56 | 5.374286  |
| 56 | 57 | 6.660000  |
| 57 | 58 | 5.517143  |
| 58 | 59 | 6.106667, |

```
Figure({
 'data': [{ 'alignmentgroup': 'True',
 'hvertemplate': 'Customer ID=%{x}
Average Invoice Total=%{y}<extra></extra>',
 'legendgroup': '',
 'marker': { 'color': '#636efa', 'pattern': { 'shape': '' } },
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,
 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36,
 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54,
 55, 56, 57, 58, 59]),
 'xaxis': 'x',
 'y': array([5.66, 5.37428571, 5.66, 5.66, 5.80285714, 7.08857143,
 6.08857143, 5.37428571, 5.37428571, 5.37428571, 5.37428571, 5.37428571,
 5.37428571, 5.37428571, 5.51714286, 5.37428571, 5.66, 5.37428571,
 5.51714286, 5.66, 5.37428571, 5.66, 5.37428571, 6.23142857,
```

```

 6.08857143, 6.80285714, 5.37428571, 6.23142857, 5.37428571, 5.37428571,
 5.37428571, 5.37428571, 5.37428571, 5.66 , 5.37428571, 5.37428571,
 6.23142857, 5.37428571, 5.51714286, 5.51714286, 5.37428571, 5.66 ,
 5.80285714, 5.94571429, 6.51714286, 6.51714286, 5.37428571, 5.80285714,
 5.37428571, 5.37428571, 5.51714286, 5.37428571, 5.37428571, 5.37428571,
 5.37428571, 5.37428571, 6.66 , 5.51714286, 6.10666667]],
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Average Invoice Total for Each Customer'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'Customer ID'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Average Invoice Total'}}}
)))

```

```

In [31]: question = """
 Find the top 5 most expensive tracks (based on unit price):
 """

 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables \nCREATE TABLE "tracks"\n\n TrackId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n Name NVARCHAR(200) NOT NULL,\n\n AlbumId INTEGER,\n\n MediaTypeId INTEGER NOT NULL,\n\n GenreId INTEGER,\n\n Composer NVARCHAR(220),\n\n Milliseconds INTEGER NOT NULL,\n\n Bytes INTEGER,\n\n UnitPrice NUMERIC(10,2) NOT NULL,\n\n FOREIGN KEY (AlbumId) REFERENCES "albums" (AlbumId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (GenreId) REFERENCES "genres" (GenreId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (MediaTypeId) REFERENCES "media\_types" (MediaTypeId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE INDEX IFK\_TrackAlbumId ON "tracks" (AlbumId)\n\nCREATE INDEX IFK\_TrackGenreId ON "tracks" (GenreId)\n\nCREATE INDEX IFK\_PlaylistTrackTrackId ON "playlist\_track" (TrackId)\n\nCREATE INDEX IFK\_InvoiceLineTrackId ON "invoice\_items" (TrackId)\n\nCREATE INDEX IFK\_TrackMediaTypeId ON "tracks" (MediaTypeId)\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n InvoiceId INTEGER NOT NULL,\n\n TrackId INTEGER NOT NULL,\n\n UnitPrice NUMERIC(10,2) NOT NULL,\n\n Quantity INTEGER NOT NULL,\n\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE TABLE "playlist\_track"\n\n PlaylistId INTEGER NOT NULL,\n\n TrackId INTEGER NOT NULL,\n\n CONSTRAINT PK\_PlaylistTrack PRIMARY KEY (PlaylistId, TrackId),\n\n FOREIGN KEY (PlaylistId) REFERENCES "playlists" (PlaylistId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\n\t\tON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\n\n===Additional Context \n\n===Response Guidelines \n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question. \n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql \n3. If the provided context is insufficient, please explain why it can't be generated. \n4. Please use the most relevant table(s). \n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before. \n'}], {'role': 'user', 'content': ' \n Find all tracks with a name containing "What" (case-insensitive)\n'}, {'role': 'assistant', 'content': "SELECT \* FROM tracks WHERE Name LIKE '%What%';"}, {'role': 'user', 'content': ' \n List all invoices with a total exceeding \$10:\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': ' \n List all albums and their corresponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': ' \n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': "SELECT \*, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n Get the average invoice total for each customer:\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal \nFROM invoices \nGROUP BY CustomerId;'}, {'role': 'user', 'content': ' \n Find the total number of invoices per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY BillingCountry;'}, {'role': 'user', 'content': ' \n Get the total number of invoices for each customer \n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY CustomerId;'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role':



'user', 'content': ' \n Find the top 5 most expensive tracks (based on unit price):\n']

Using model gpt-4o-mini for 981.0 tokens (approx)

LLM Response: SELECT \* FROM tracks ORDER BY UnitPrice DESC LIMIT 5;

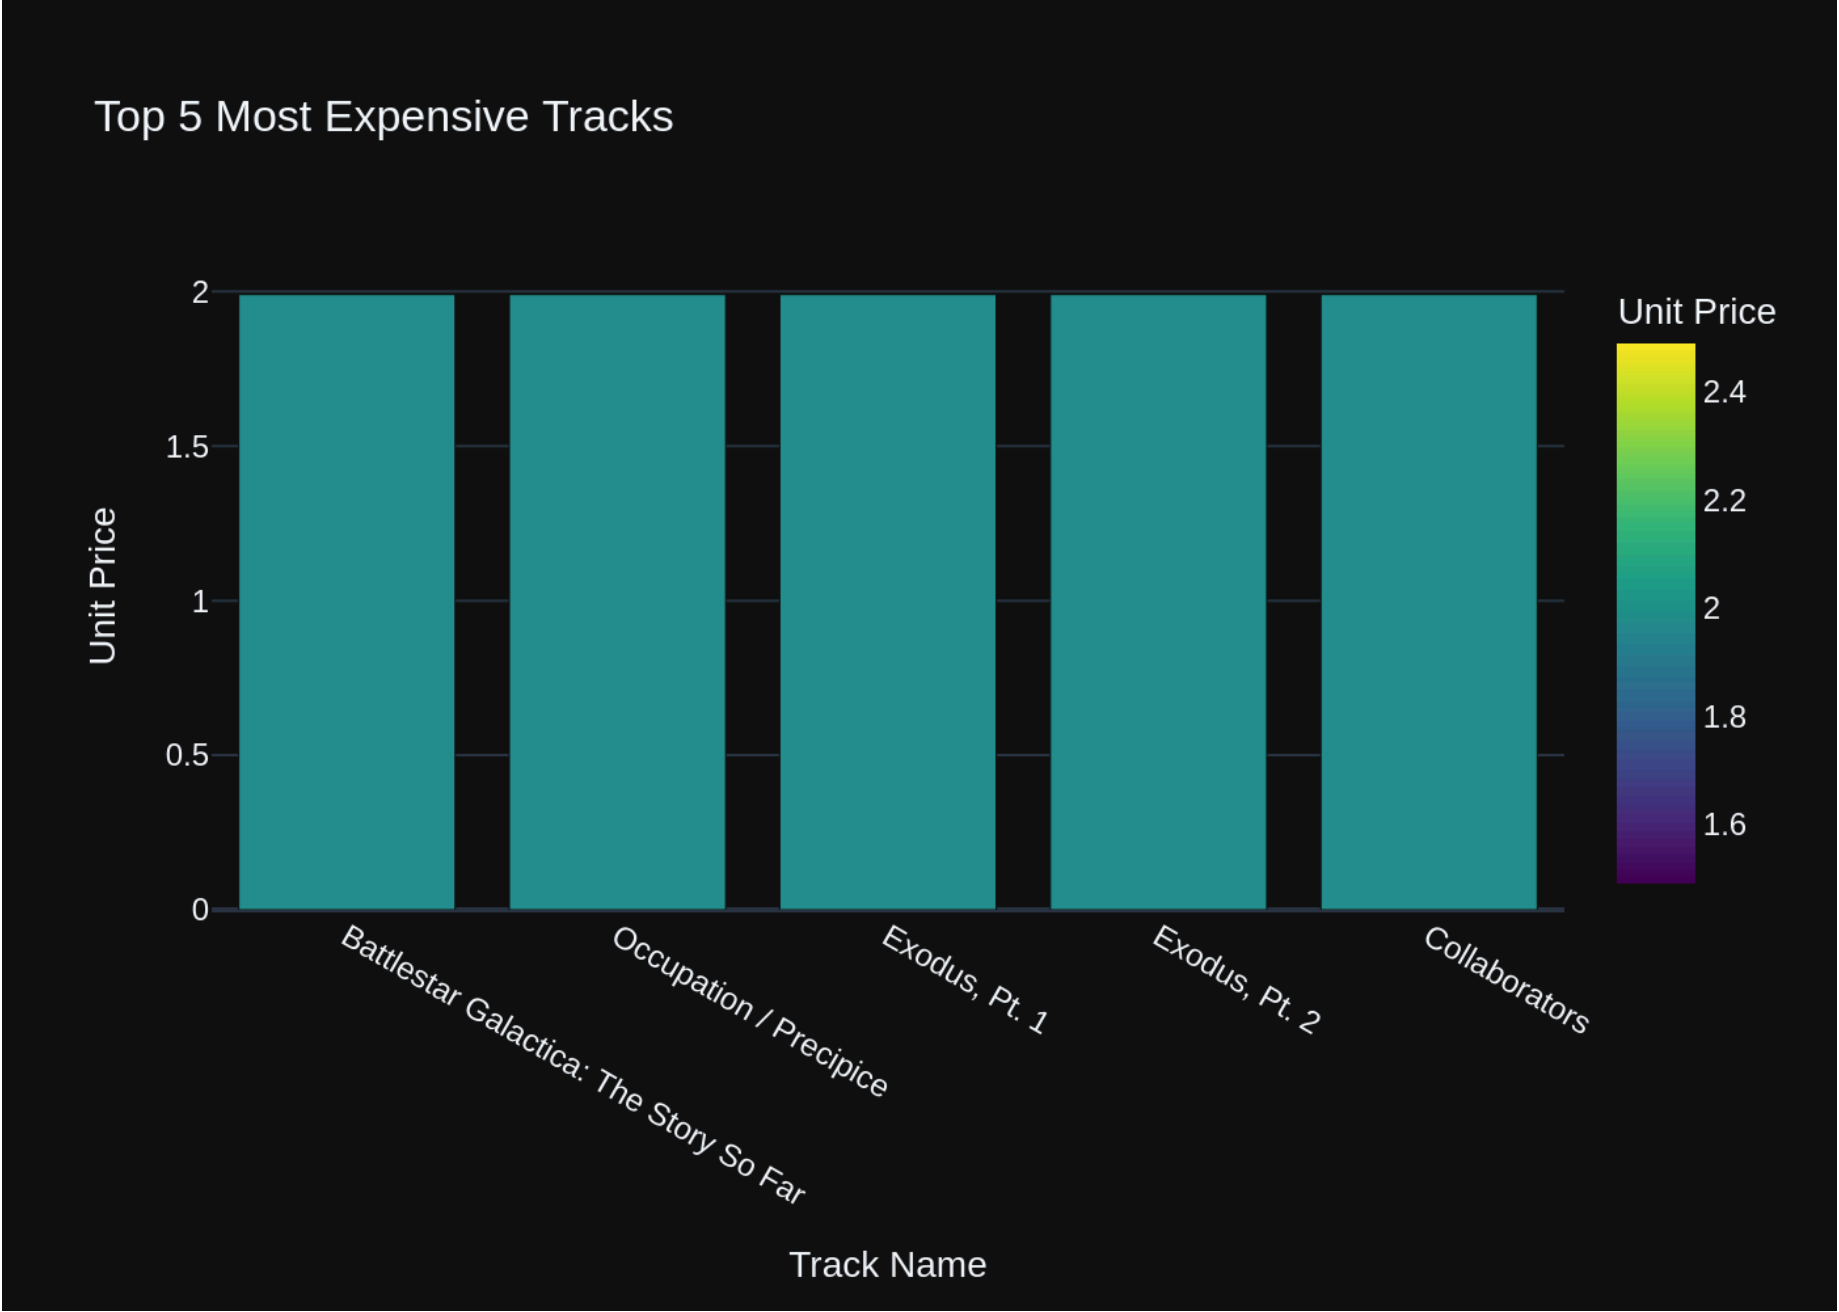
Extracted SQL: SELECT \* FROM tracks ORDER BY UnitPrice DESC LIMIT 5;

SELECT \* FROM tracks ORDER BY UnitPrice DESC LIMIT 5;

|   | TrackId | Name                                   | AlbumId | MediaTypeId | \ |
|---|---------|----------------------------------------|---------|-------------|---|
| 0 | 2819    | Battlestar Galactica: The Story So Far | 226     | 3           |   |
| 1 | 2820    | Occupation / Precipice                 | 227     | 3           |   |
| 2 | 2821    | Exodus, Pt. 1                          | 227     | 3           |   |
| 3 | 2822    | Exodus, Pt. 2                          | 227     | 3           |   |
| 4 | 2823    | Collaborators                          | 227     | 3           |   |

|   | GenreId | Composer | Milliseconds | Bytes      | UnitPrice |
|---|---------|----------|--------------|------------|-----------|
| 0 | 18      | None     | 2622250      | 490750393  | 1.99      |
| 1 | 19      | None     | 5286953      | 1054423946 | 1.99      |
| 2 | 19      | None     | 2621708      | 475079441  | 1.99      |
| 3 | 19      | None     | 2618000      | 466820021  | 1.99      |
| 4 | 19      | None     | 2626626      | 483484911  | 1.99      |

Using model gpt-4o-mini for 224.0 tokens (approx)



```
Out[31]: ('SELECT * FROM tracks ORDER BY UnitPrice DESC LIMIT 5;',
```

|   | TrackId | Name                                   | AlbumId | MediaTypeId | \ |
|---|---------|----------------------------------------|---------|-------------|---|
| 0 | 2819    | Battlestar Galactica: The Story So Far | 226     | 3           |   |
| 1 | 2820    | Occupation / Precipice                 | 227     | 3           |   |
| 2 | 2821    | Exodus, Pt. 1                          | 227     | 3           |   |
| 3 | 2822    | Exodus, Pt. 2                          | 227     | 3           |   |
| 4 | 2823    | Collaborators                          | 227     | 3           |   |

|   | GenreId | Composer | Milliseconds | Bytes      | UnitPrice |
|---|---------|----------|--------------|------------|-----------|
| 0 | 18      | None     | 2622250      | 490750393  | 1.99      |
| 1 | 19      | None     | 5286953      | 1054423946 | 1.99      |
| 2 | 19      | None     | 2621708      | 475079441  | 1.99      |
| 3 | 19      | None     | 2618000      | 466820021  | 1.99      |
| 4 | 19      | None     | 2626626      | 483484911  | 1.99      |

```
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovernplate': 'Track Name=%{x}
Unit Price=%{marker.color}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': array([1.99, 1.99, 1.99, 1.99, 1.99]),
 'coloraxis': 'coloraxis',
 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array(['Battlestar Galactica: The Story So Far', 'Occupation / Precipice',
 'Exodus, Pt. 1', 'Exodus, Pt. 2', 'Collaborators'], dtype=object),
 'xaxis': 'x',
 'y': array([1.99, 1.99, 1.99, 1.99, 1.99]),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'coloraxis': {'colorbar': {'title': {'text': 'Unit Price'}}},
 'colorscale': [[0.0, '#440154'], [0.11111111111111111,
 '#482878'], [0.22222222222222222,
 '#3e4989'], [0.33333333333333333,
 '#31688e'], [0.44444444444444444,
 '#26828e'], [0.55555555555555556,
 '#1f9e89'], [0.66666666666666666,
```

```

 '#35b779'], [0.7777777777777778,
 '#6ece58'], [0.8888888888888888,
 '#b5de2b'], [1.0, '#fde725']]],
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Top 5 Most Expensive Tracks'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'Track Name'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Unit Price'}}}
)))

```

```

In [32]: question = """
 List all genres and the number of tracks in each genre:
 """

 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

file:///home/gongai/Downloads/openai-gpt-4o-mini-chromadb-sqlite-test-1.html

genres and the number of tracks in each genre:\n'}]

Using model gpt-4o-mini for 966.75 tokens (approx)

LLM Response: SELECT genres.Name AS GenreName, COUNT(tracks.TrackId) AS TrackCount  
FROM genres

LEFT JOIN tracks ON genres.GenreId = tracks.GenreId

GROUP BY genres.GenreId;

Extracted SQL: SELECT genres.Name AS GenreName, COUNT(tracks.TrackId) AS TrackCount  
FROM genres

LEFT JOIN tracks ON genres.GenreId = tracks.GenreId

GROUP BY genres.GenreId;


SELECT genres.Name AS GenreName, COUNT(tracks.TrackId) AS TrackCount  
FROM genres

LEFT JOIN tracks ON genres.GenreId = tracks.GenreId

GROUP BY genres.GenreId;

|    | GenreName          | TrackCount |
|----|--------------------|------------|
| 0  | Rock               | 1297       |
| 1  | Jazz               | 130        |
| 2  | Metal              | 374        |
| 3  | Alternative & Punk | 332        |
| 4  | Rock And Roll      | 12         |
| 5  | Blues              | 81         |
| 6  | Latin              | 579        |
| 7  | Reggae             | 58         |
| 8  | Pop                | 48         |
| 9  | Soundtrack         | 43         |
| 10 | Bossa Nova         | 15         |
| 11 | Easy Listening     | 24         |
| 12 | Heavy Metal        | 28         |
| 13 | R&B/Soul           | 61         |
| 14 | Electronica/Dance  | 30         |
| 15 | World              | 28         |
| 16 | Hip Hop/Rap        | 35         |
| 17 | Science Fiction    | 13         |
| 18 | TV Shows           | 93         |
| 19 | Sci Fi & Fantasy   | 26         |
| 20 | Drama              | 64         |
| 21 | Comedy             | 17         |
| 22 | Alternative        | 40         |
| 23 | Classical          | 74         |
| 24 | Opera              | 1          |

Using model gpt-4o-mini for 206.25 tokens (approx)



```
Out[32]: ('SELECT genres.Name AS GenreName, COUNT(tracks.TrackId) AS TrackCount \nFROM genres \nLEFT JOIN tracks ON genres.\nGenreId = tracks.GenreId \nGROUP BY genres.GenreId;',
```

|    | GenreName          | TrackCount |
|----|--------------------|------------|
| 0  | Rock               | 1297       |
| 1  | Jazz               | 130        |
| 2  | Metal              | 374        |
| 3  | Alternative & Punk | 332        |
| 4  | Rock And Roll      | 12         |
| 5  | Blues              | 81         |
| 6  | Latin              | 579        |
| 7  | Reggae             | 58         |
| 8  | Pop                | 48         |
| 9  | Soundtrack         | 43         |
| 10 | Bossa Nova         | 15         |
| 11 | Easy Listening     | 24         |
| 12 | Heavy Metal        | 28         |
| 13 | R&B/Soul           | 61         |
| 14 | Electronica/Dance  | 30         |
| 15 | World              | 28         |
| 16 | Hip Hop/Rap        | 35         |
| 17 | Science Fiction    | 13         |
| 18 | TV Shows           | 93         |
| 19 | Sci Fi & Fantasy   | 26         |
| 20 | Drama              | 64         |
| 21 | Comedy             | 17         |
| 22 | Alternative        | 40         |
| 23 | Classical          | 74         |
| 24 | Opera              | 1,         |

```
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovernplate': 'GenreName=%{x}
TrackCount=%{y}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': '#636efa', 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array(['Rock', 'Jazz', 'Metal', 'Alternative & Punk', 'Rock And Roll', 'Blues',
```



```

 'Latin', 'Reggae', 'Pop', 'Soundtrack', 'Bossa Nova', 'Easy Listening',
 'Heavy Metal', 'R&B/Soul', 'Electronica/Dance', 'World', 'Hip Hop/Rap',
 'Science Fiction', 'TV Shows', 'Sci Fi & Fantasy', 'Drama', 'Comedy',
 'Alternative', 'Classical', 'Opera'], dtype=object),
 'xaxis': 'x',
 'y': array([1297, 130, 374, 332, 12, 81, 579, 58, 48, 43, 15, 24,
 28, 61, 30, 28, 35, 13, 93, 26, 64, 17, 40, 74,
 1]),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Number of Tracks per Genre'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'GenreName'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'TrackCount'}}}
)))

```

```

In [33]: question = """
 Get all genres that do not have any tracks associated with them:
 """

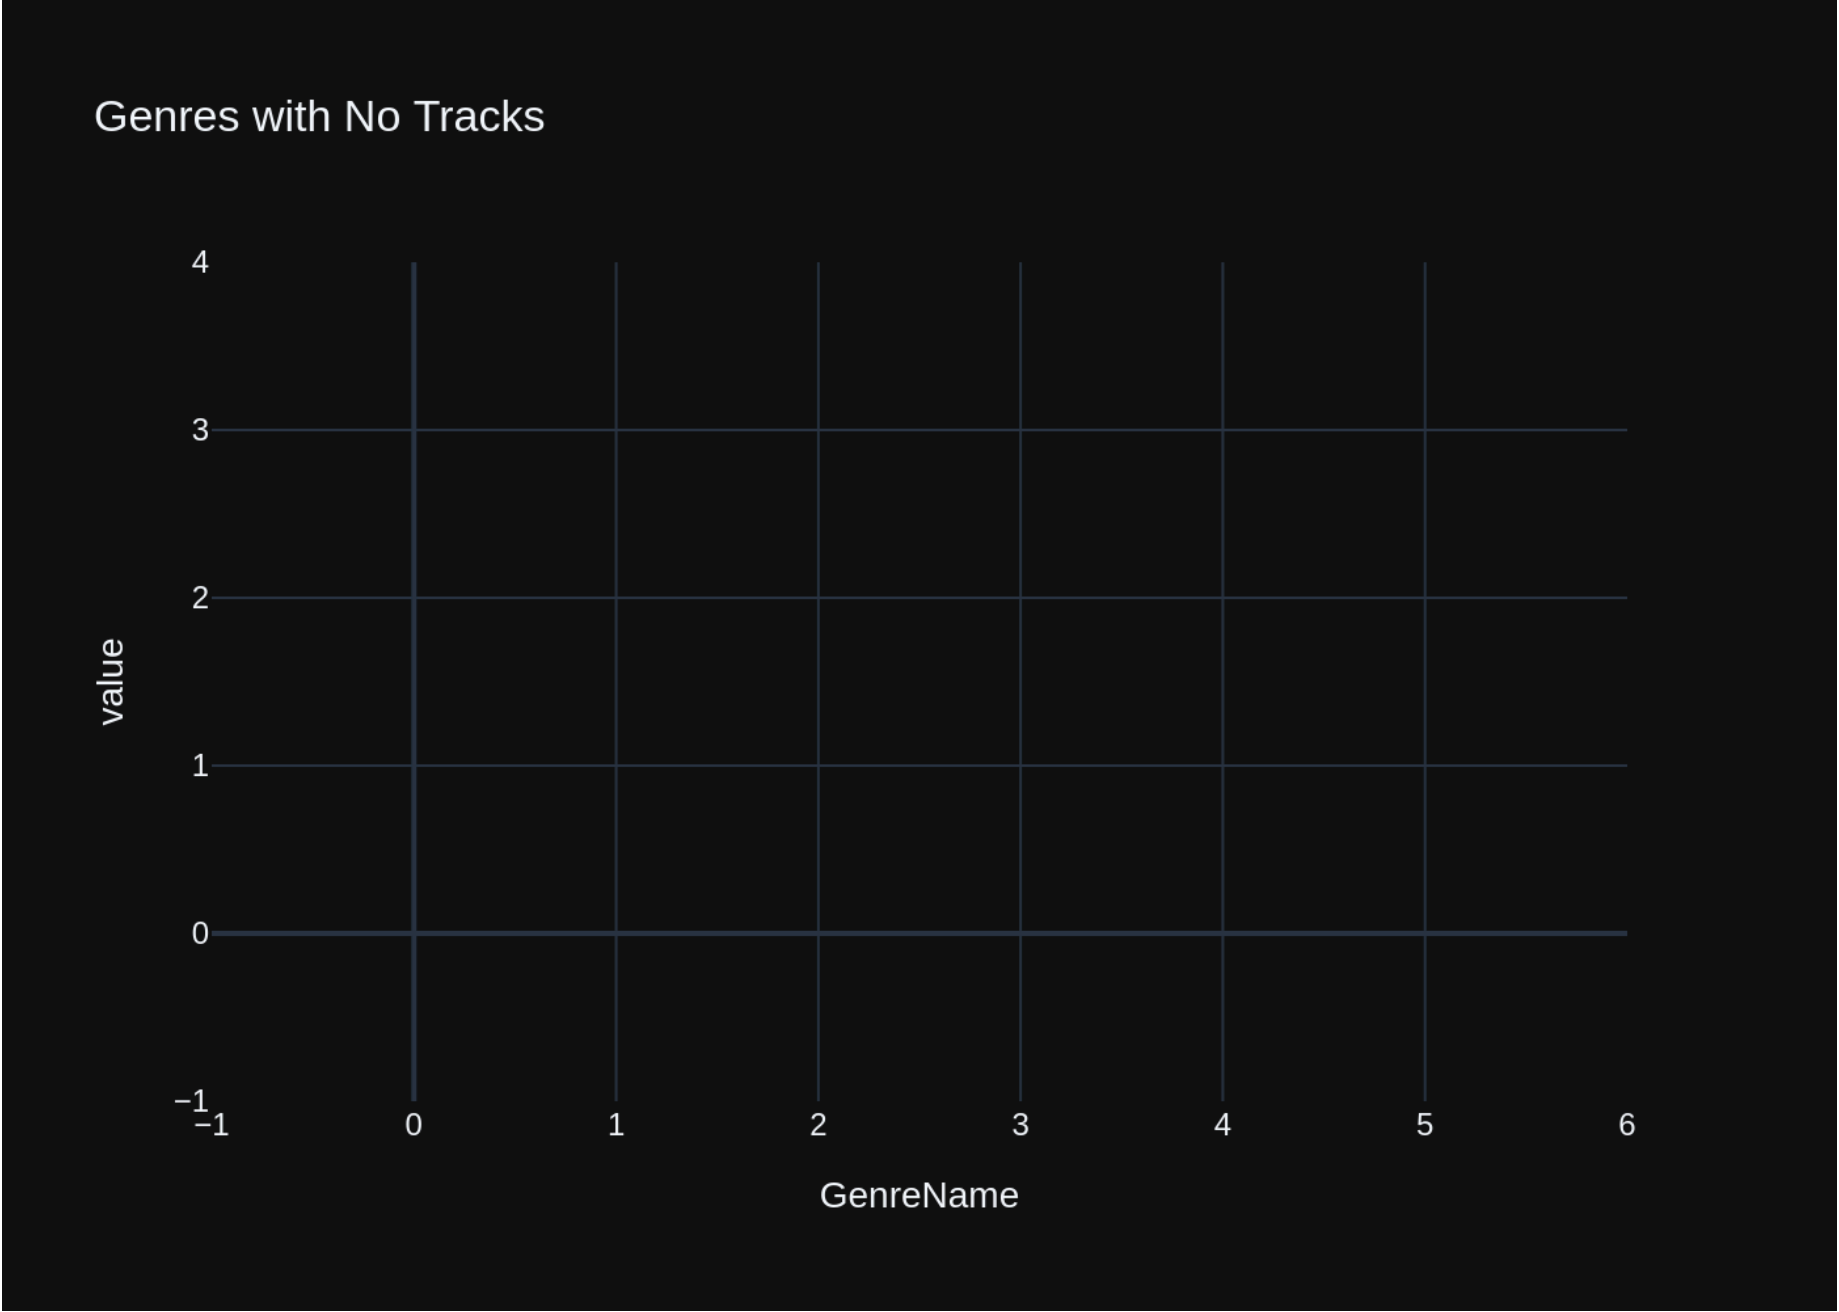
 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables \nCREATE INDEX IFK\_TrackGenreId ON "tracks" (GenreId)\n\nCREATE TABLE "tracks"\n\nTrackId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\nName NVARCHAR(200) NOT NULL,\n\nAlbumId INTEGER,\n\nMediaTypeId INTEGER NOT NULL,\n\nGenreId INTEGER,\n\nComposer NVARCHAR(220),\n\nMilliseconds INTEGER NOT NULL,\n\nBytes INTEGER,\n\nUnitPrice NUMERIC(10,2) NOT NULL,\n\nFOREIGN KEY (AlbumId) REFERENCES "albums" (AlbumId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\nFOREIGN KEY (GenreId) REFERENCES "genres" (GenreId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\nFOREIGN KEY (MediaTypeId) REFERENCES "media\_types" (MediaTypeId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE INDEX IFK\_PlaylistTrackTrackId ON "playlist\_track" (TrackId)\n\nCREATE INDEX IFK\_TrackMediaTypeId ON "tracks" (MediaTypeId)\n\nCREATE INDEX IFK\_TrackAlbumId ON "tracks" (AlbumId)\n\nCREATE TABLE "genres"\n\nGenreId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\nName NVARCHAR(120)\n\n)\n\nCREATE TABLE "albums"\n\nAlbumId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\nTitle NVARCHAR(160) NOT NULL,\n\nArtistId INTEGER NOT NULL,\n\nFOREIGN KEY (ArtistId) REFERENCES "artists" (ArtistId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE TABLE "playlist\_track"\n\nPlaylistId INTEGER NOT NULL,\n\nTrackId INTEGER NOT NULL,\n\nCONSTRAINT PK\_PlaylistTrack PRIMARY KEY (PlaylistId, TrackId),\n\nFOREIGN KEY (PlaylistId) REFERENCES "playlists" (PlaylistId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\nFOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n)\n\nCREATE INDEX IFK\_AlbumArtistId ON "albums" (ArtistId)\n\n\n===Additional Context \n\nIn the chinook database invoice means order\n\n===Response Guidelines \n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question. \n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql \n3. If the provided context is insufficient, please explain why it can't be generated. \n4. Please use the most relevant table(s). \n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before. \n'}], {'role': 'user', 'content': ' \n List all genres and the number of tracks in each genre:\n'}, {'role': 'assistant', 'content': 'SELECT genres.Name AS GenreName, COUNT(tracks.TrackId) AS TrackCount \nFROM genres \nLEFT JOIN tracks ON genres.GenreId = tracks.GenreId \nGROUP BY genres.GenreId;'}, {'role': 'user', 'content': ' \n Find all tracks with a name containing "What" (case-insensitive)\n'}, {'role': 'assistant', 'content': "SELECT \* FROM tracks WHERE Name LIKE '%What%';"}, {'role': 'user', 'content': ' \n List all albums and their corresponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': ' \n Find the top 5 most expensive tracks (based on unit price):\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM tracks ORDER BY UnitPrice DESC LIMIT 5;'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': ' \n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': "SELECT \*, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n List all invoices with a total exceeding \$10:\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': ' \n List all employees and their reporting manager's name (if any):\n'}, {'role': 'assistant', 'content': 'SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName, \n m.FirstName AS ManagerFirstName, m.LastName AS ManagerLastName\nFROM employees e\nLEFT JOIN employees m ON e.ReportsTo = m.E'}]

```
mployeeId;'}, {'role': 'user', 'content': ' \n Find the total number of invoices per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(*) AS TotalInvoices \nFROM invoices \nGROUP BY BillingCountry;'}, {'role': 'user', 'content': ' \n Get all genres that do not have any tracks associated with them:\n'}]
Using model gpt-4o-mini for 1041.5 tokens (approx)
LLM Response: SELECT genres.Name AS GenreName
FROM genres
LEFT JOIN tracks ON genres.GenreId = tracks.GenreId
WHERE tracks.TrackId IS NULL;
Extracted SQL: SELECT genres.Name AS GenreName
FROM genres
LEFT JOIN tracks ON genres.GenreId = tracks.GenreId
WHERE tracks.TrackId IS NULL;
SELECT genres.Name AS GenreName
FROM genres
LEFT JOIN tracks ON genres.GenreId = tracks.GenreId
WHERE tracks.TrackId IS NULL;
Empty DataFrame
Columns: [GenreName]
Index: []
Using model gpt-4o-mini for 195.0 tokens (approx)
```



```

Out[33]: ('SELECT genres.Name AS GenreName \nFROM genres \nLEFT JOIN tracks ON genres.GenreId = tracks.GenreId \nWHERE trac
ks.TrackId IS NULL;',
Empty DataFrame
Columns: [GenreName]
Index: [],
Figure({
 'data': [],
 'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Genres with No Tracks'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'GenreName'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'value'}}}
}))

```

```

In [34]: question = """
List all customers who have not placed any orders:
"""

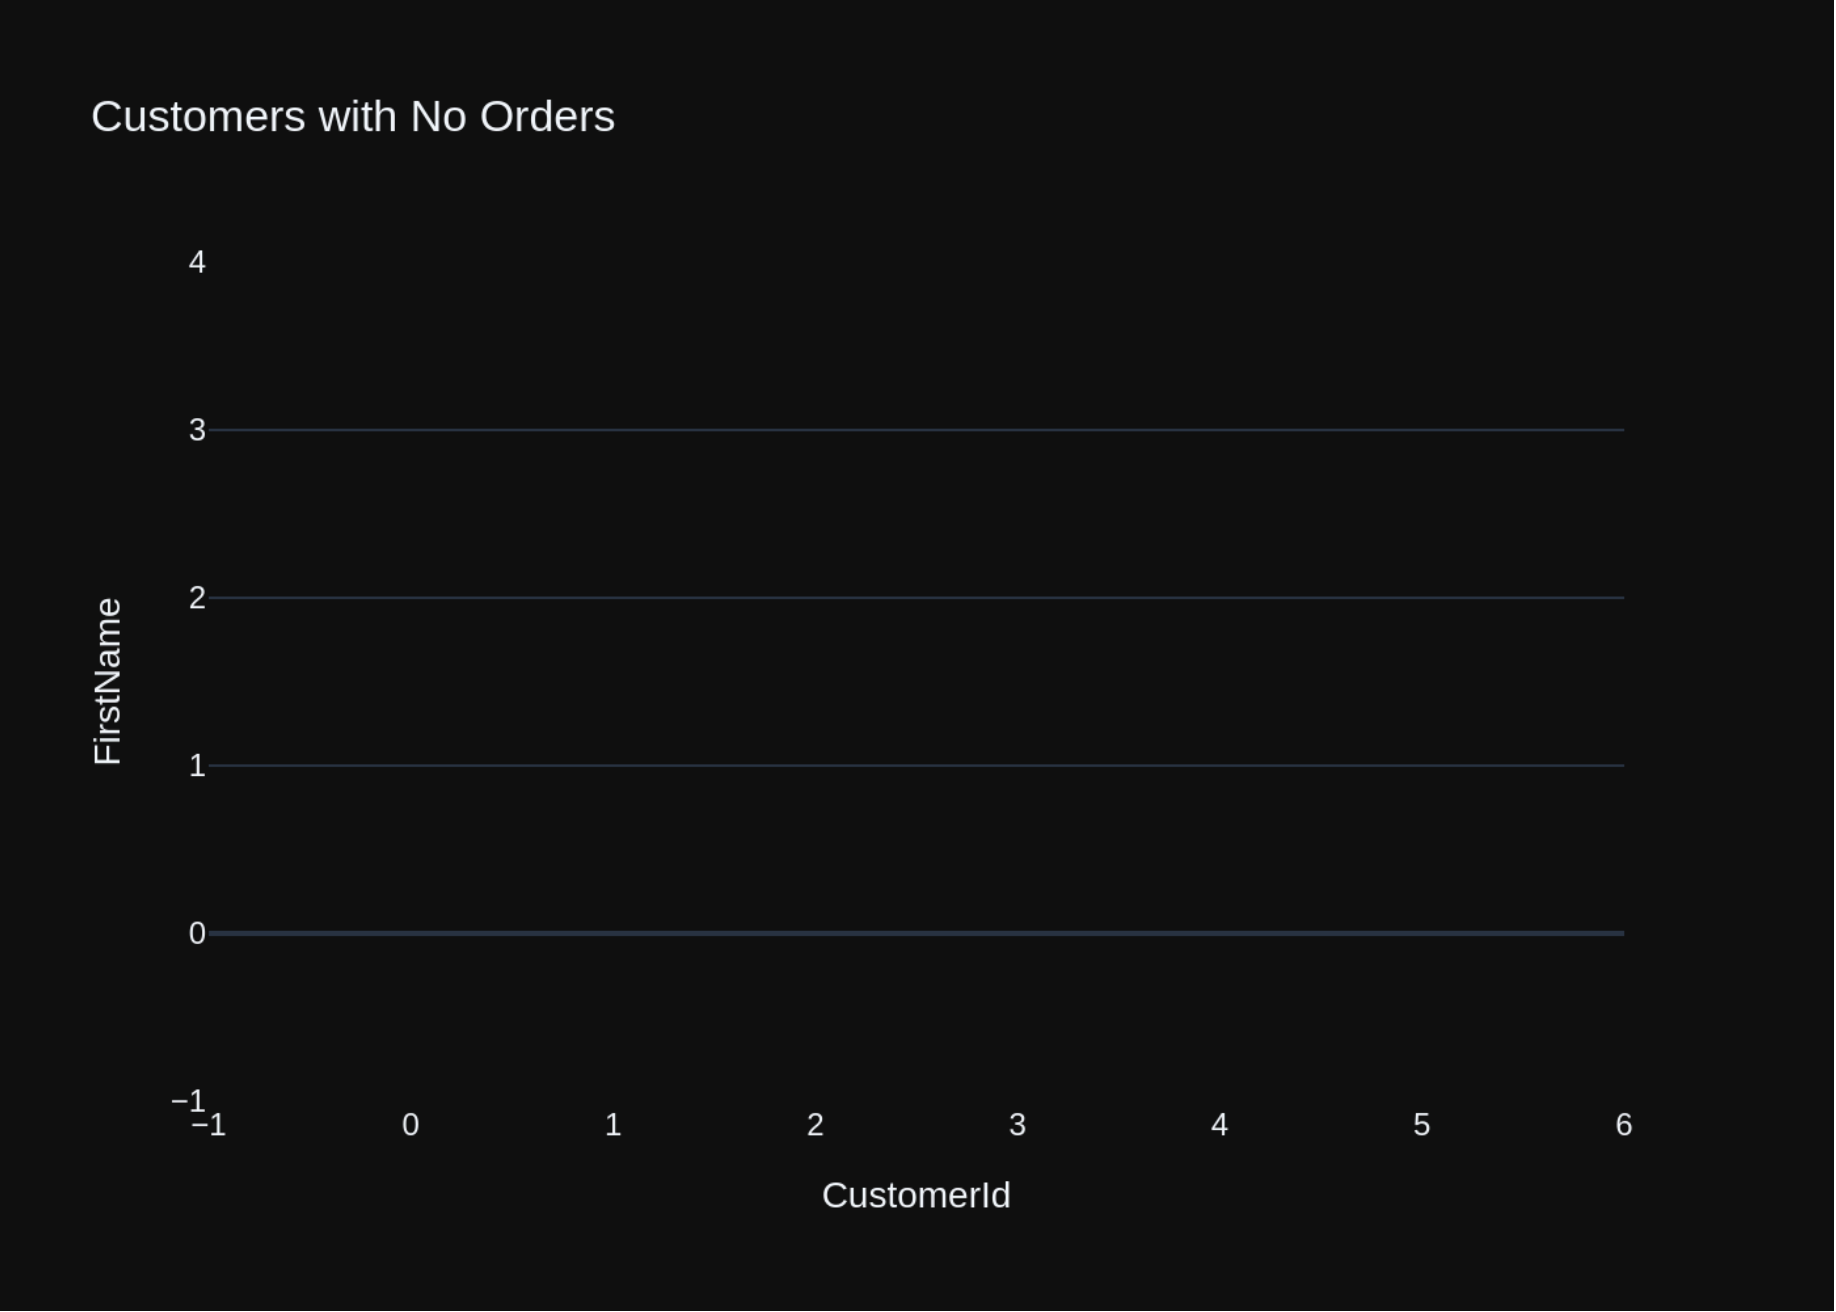
vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n CustomerId INTEGER NOT NULL,\n\n InvoiceDate DATETIME NOT NULL,\n\n BillingAddress NVARCHAR(70),\n\n BillingCity NVARCHAR(40),\n\n BillingState NVARCHAR(40),\n\n BillingCountry NVARCHAR(40),\n\n BillingPostalCode NVARCHAR(10),\n\n Total NUMERIC(10,2) NOT NULL,\n\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE TABLE "customers"\n\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n FirstName NVARCHAR(40) NOT NULL,\n\n LastName NVARCHAR(20) NOT NULL,\n\n Company NVARCHAR(80),\n\n Address NVARCHAR(70),\n\n City NVARCHAR(40),\n\n State NVARCHAR(40),\n\n Country NVARCHAR(40),\n\n PostalCode NVARCHAR(10),\n\n Phone NVARCHAR(24),\n\n Fax NVARCHAR(24),\n\n Email NVARCHAR(60) NOT NULL,\n\n SupportRepId INTEGER,\n\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n InvoiceId INTEGER NOT NULL,\n\n TrackId INTEGER NOT NULL,\n\n UnitPrice NUMERIC(10,2) NOT NULL,\n\n Quantity INTEGER NOT NULL,\n\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\nCREATE TABLE "albums"\n\n AlbumId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n\n Title NVARCHAR(160) NOT NULL,\n\n ArtistId INTEGER NOT NULL,\n\n FOREIGN KEY (ArtistId) REFERENCES "artists" (ArtistId) \n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\n\n\n===Additional Context\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'}], {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': ' \n Get the total number of invoices for each customer\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY CustomerId;'}, {'role': 'user', 'content': ' \n List all invoices with a total exceeding \$10:\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': ' \n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': 'SELECT \*, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n Get the average invoice total for each customer:\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal \nFROM invoices \nGROUP BY CustomerId;'}, {'role': 'user', 'content': ' \n Find the total number of invoices per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(\*) AS TotalInvoices \nFROM invoices \nGROUP BY BillingCountry;'}, {'role': 'user', 'content': ' \n List all employees and their reporting manager's name (if any):\n'}, {'role': 'assistant', 'content': 'SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName, \n m.FirstName AS ManagerFirstName, m.LastName AS ManagerLastName\nFROM employees e\nLEFT JOIN employees m ON e.ReportsTo = m.EmployeeId;'}, {'role': 'user', 'content': ' \n List all albums and their corresponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user',

```
'content': ' \n Find the top 5 most expensive tracks (based on unit price):\n', {'role': 'assistant', 'content': 'SELECT * FROM tracks ORDER BY UnitPrice DESC LIMIT 5;'}, {'role': 'user', 'content': ' \n List all customers who have not placed any orders:\n'}]
Using model gpt-4o-mini for 1018.0 tokens (approx)
LLM Response: SELECT * FROM customers
WHERE CustomerId NOT IN (SELECT DISTINCT CustomerId FROM invoices);
Extracted SQL: SELECT * FROM customers
WHERE CustomerId NOT IN (SELECT DISTINCT CustomerId FROM invoices);
SELECT * FROM customers
WHERE CustomerId NOT IN (SELECT DISTINCT CustomerId FROM invoices);
Empty DataFrame
Columns: [CustomerId, FirstName, LastName, Company, Address, City, State, Country, PostalCode, Phone, Fax, Email, SupportRepId]
Index: []
Using model gpt-4o-mini for 252.25 tokens (approx)
```





```

Out[34]: ('SELECT * FROM customers \nWHERE CustomerId NOT IN (SELECT DISTINCT CustomerId FROM invoices);',
Empty DataFrame
Columns: [CustomerId, FirstName, LastName, Company, Address, City, State, Country, PostalCode, Phone, Fax, Email,
SupportRepId]
Index: [],
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovernment': 'CustomerId=%{x}
FirstName=%{y}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': '#636efa', 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array([], dtype=object),
 'xaxis': 'x',
 'y': array([], dtype=object),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Customers with No Orders'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'CustomerId'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'FirstName'}}
}))

```

```

In [35]: question = """
 Get the top 10 most popular artists (based on the number of tracks):
 """

vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

90/129

Using model gpt-4o-mini for 933.75 tokens (approx)

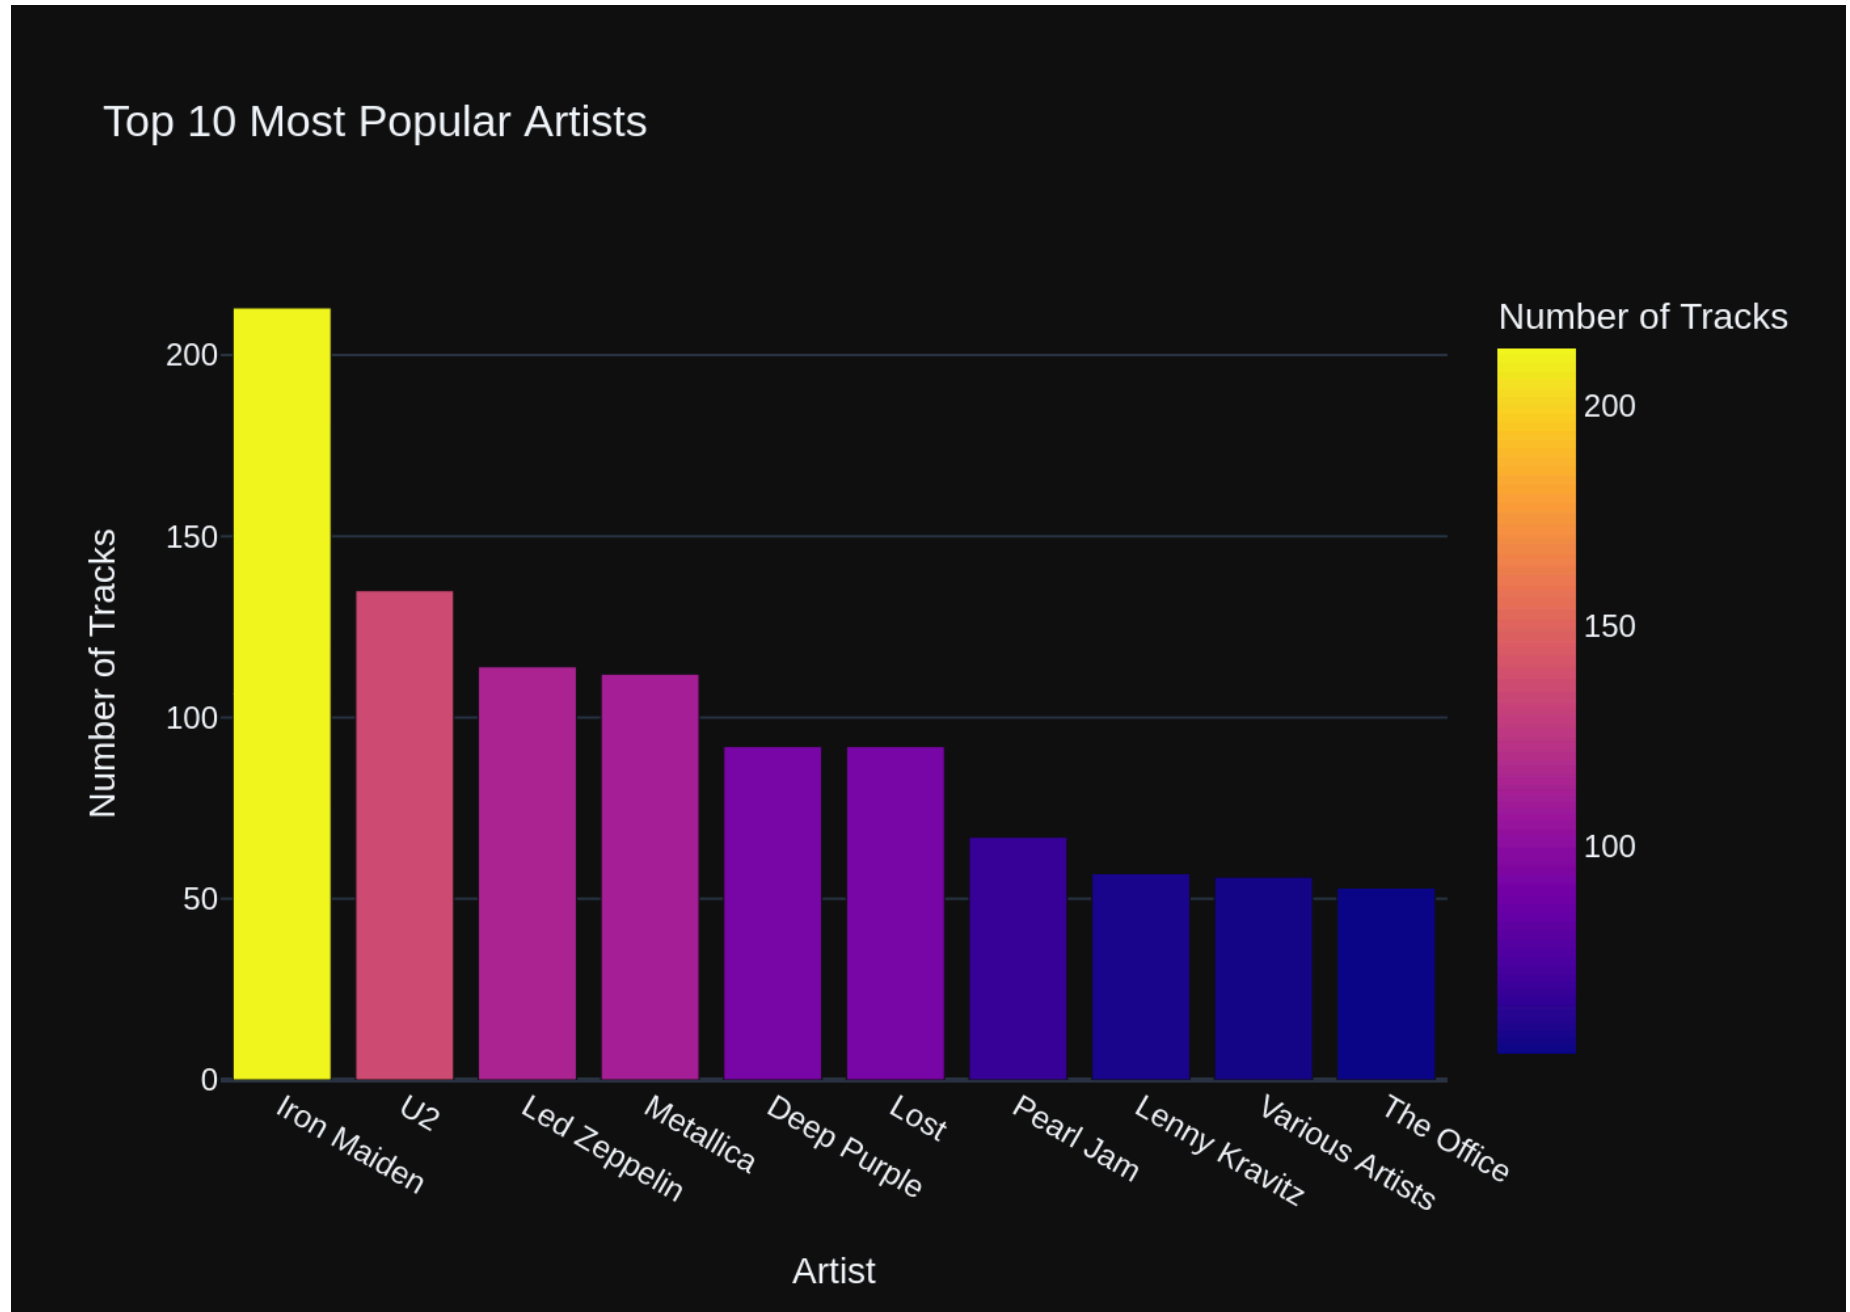
```
LLM Response: SELECT artists.Name AS ArtistName, COUNT(tracks.TrackId) AS TrackCount
FROM artists
JOIN albums ON artists.ArtistId = albums.ArtistId
JOIN tracks ON albums.AlbumId = tracks.AlbumId
GROUP BY artists.ArtistId
ORDER BY TrackCount DESC
LIMIT 10;
```

```
Extracted SQL: SELECT artists.Name AS ArtistName, COUNT(tracks.TrackId) AS TrackCount
FROM artists
JOIN albums ON artists.ArtistId = albums.ArtistId
JOIN tracks ON albums.AlbumId = tracks.AlbumId
GROUP BY artists.ArtistId
ORDER BY TrackCount DESC
LIMIT 10;
```

```
SELECT artists.Name AS ArtistName, COUNT(tracks.TrackId) AS TrackCount
FROM artists
JOIN albums ON artists.ArtistId = albums.ArtistId
JOIN tracks ON albums.AlbumId = tracks.AlbumId
GROUP BY artists.ArtistId
ORDER BY TrackCount DESC
LIMIT 10;
```

|   | ArtistName      | TrackCount |
|---|-----------------|------------|
| 0 | Iron Maiden     | 213        |
| 1 | U2              | 135        |
| 2 | Led Zeppelin    | 114        |
| 3 | Metallica       | 112        |
| 4 | Deep Purple     | 92         |
| 5 | Lost            | 92         |
| 6 | Pearl Jam       | 67         |
| 7 | Lenny Kravitz   | 57         |
| 8 | Various Artists | 56         |
| 9 | The Office      | 53         |

Using model gpt-4o-mini for 231.25 tokens (approx)



```
Out[35]: ('SELECT artists.Name AS ArtistName, COUNT(tracks.TrackId) AS TrackCount \nFROM artists \nJOIN albums ON artists.A\nrtistId = albums.ArtistId \nJOIN tracks ON albums.AlbumId = tracks.AlbumId \nGROUP BY artists.ArtistId \nORDER BY\nTrackCount DESC \nLIMIT 10;',
```

|   | ArtistName      | TrackCount |
|---|-----------------|------------|
| 0 | Iron Maiden     | 213        |
| 1 | U2              | 135        |
| 2 | Led Zeppelin    | 114        |
| 3 | Metallica       | 112        |
| 4 | Deep Purple     | 92         |
| 5 | Lost            | 92         |
| 6 | Pearl Jam       | 67         |
| 7 | Lenny Kravitz   | 57         |
| 8 | Various Artists | 56         |
| 9 | The Office      | 53,        |

```
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovernment': 'Artist=%{x}
Number of Tracks=%{marker.color}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': array([213, 135, 114, 112, 92, 92, 67, 57, 56, 53]),
 'coloraxis': 'coloraxis',
 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array(['Iron Maiden', 'U2', 'Led Zeppelin', 'Metallica', 'Deep Purple', 'Lost',
 'Pearl Jam', 'Lenny Kravitz', 'Various Artists', 'The Office'],
 dtype=object),
 'xaxis': 'x',
 'y': array([213, 135, 114, 112, 92, 92, 67, 57, 56, 53]),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'coloraxis': {'colorbar': {'title': {'text': 'Number of Tracks'}}},
 'colorscale': [[0.0, '#0d0887'], [0.1111111111111111,
 '#46039f'], [0.2222222222222222,
 '#7201a8'], [0.3333333333333333,
 '#9c179e'], [0.4444444444444444,
 '#bd3786'], [0.5555555555555555,
```

```

 '#d8576b'], [0.6666666666666666,
 '#ed7953'], [0.7777777777777778,
 '#fb9f3a'], [0.8888888888888888,
 '#fdca26'], [1.0, '#f0f921']]],
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Top 10 Most Popular Artists'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'Artist'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Number of Tracks'}}}
)))

```

```

In [36]: question = """
 List all customers from Canada and their email addresses:
 """

 vn.ask(question=question)

```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

95/129

```
\nJOIN albums ON artists.ArtistId = albums.ArtistId \nJOIN tracks ON albums.AlbumId = tracks.AlbumId \nGROUP BY arti
sts.ArtistId \nORDER BY TrackCount DESC \nLIMIT 10;'}}, {'role': 'user', 'content': ' \n List all albums and thei
r corresponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Nam
e AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': '
\n List all customers from Canada and their email addresses:\n'}]
```

Using model gpt-4o-mini for 1082.75 tokens (approx)

LLM Response: SELECT FirstName, LastName, Email

FROM customers

WHERE Country = 'Canada';

Extracted SQL: SELECT FirstName, LastName, Email

FROM customers

WHERE Country = 'Canada';

SELECT FirstName, LastName, Email

FROM customers

WHERE Country = 'Canada';

|  | FirstName | LastName | Email |
|--|-----------|----------|-------|
|--|-----------|----------|-------|

|   |          |          |                     |
|---|----------|----------|---------------------|
| 0 | François | Tremblay | ftremblay@gmail.com |
|---|----------|----------|---------------------|

|   |      |         |                    |
|---|------|---------|--------------------|
| 1 | Mark | Philips | mphilips12@shaw.ca |
|---|------|---------|--------------------|

|   |          |          |                     |
|---|----------|----------|---------------------|
| 2 | Jennifer | Peterson | jenniferp@rogers.ca |
|---|----------|----------|---------------------|

|   |        |       |                  |
|---|--------|-------|------------------|
| 3 | Robert | Brown | robbrown@shaw.ca |
|---|--------|-------|------------------|

|   |        |         |                     |
|---|--------|---------|---------------------|
| 4 | Edward | Francis | edfrancis@yachoo.ca |
|---|--------|---------|---------------------|

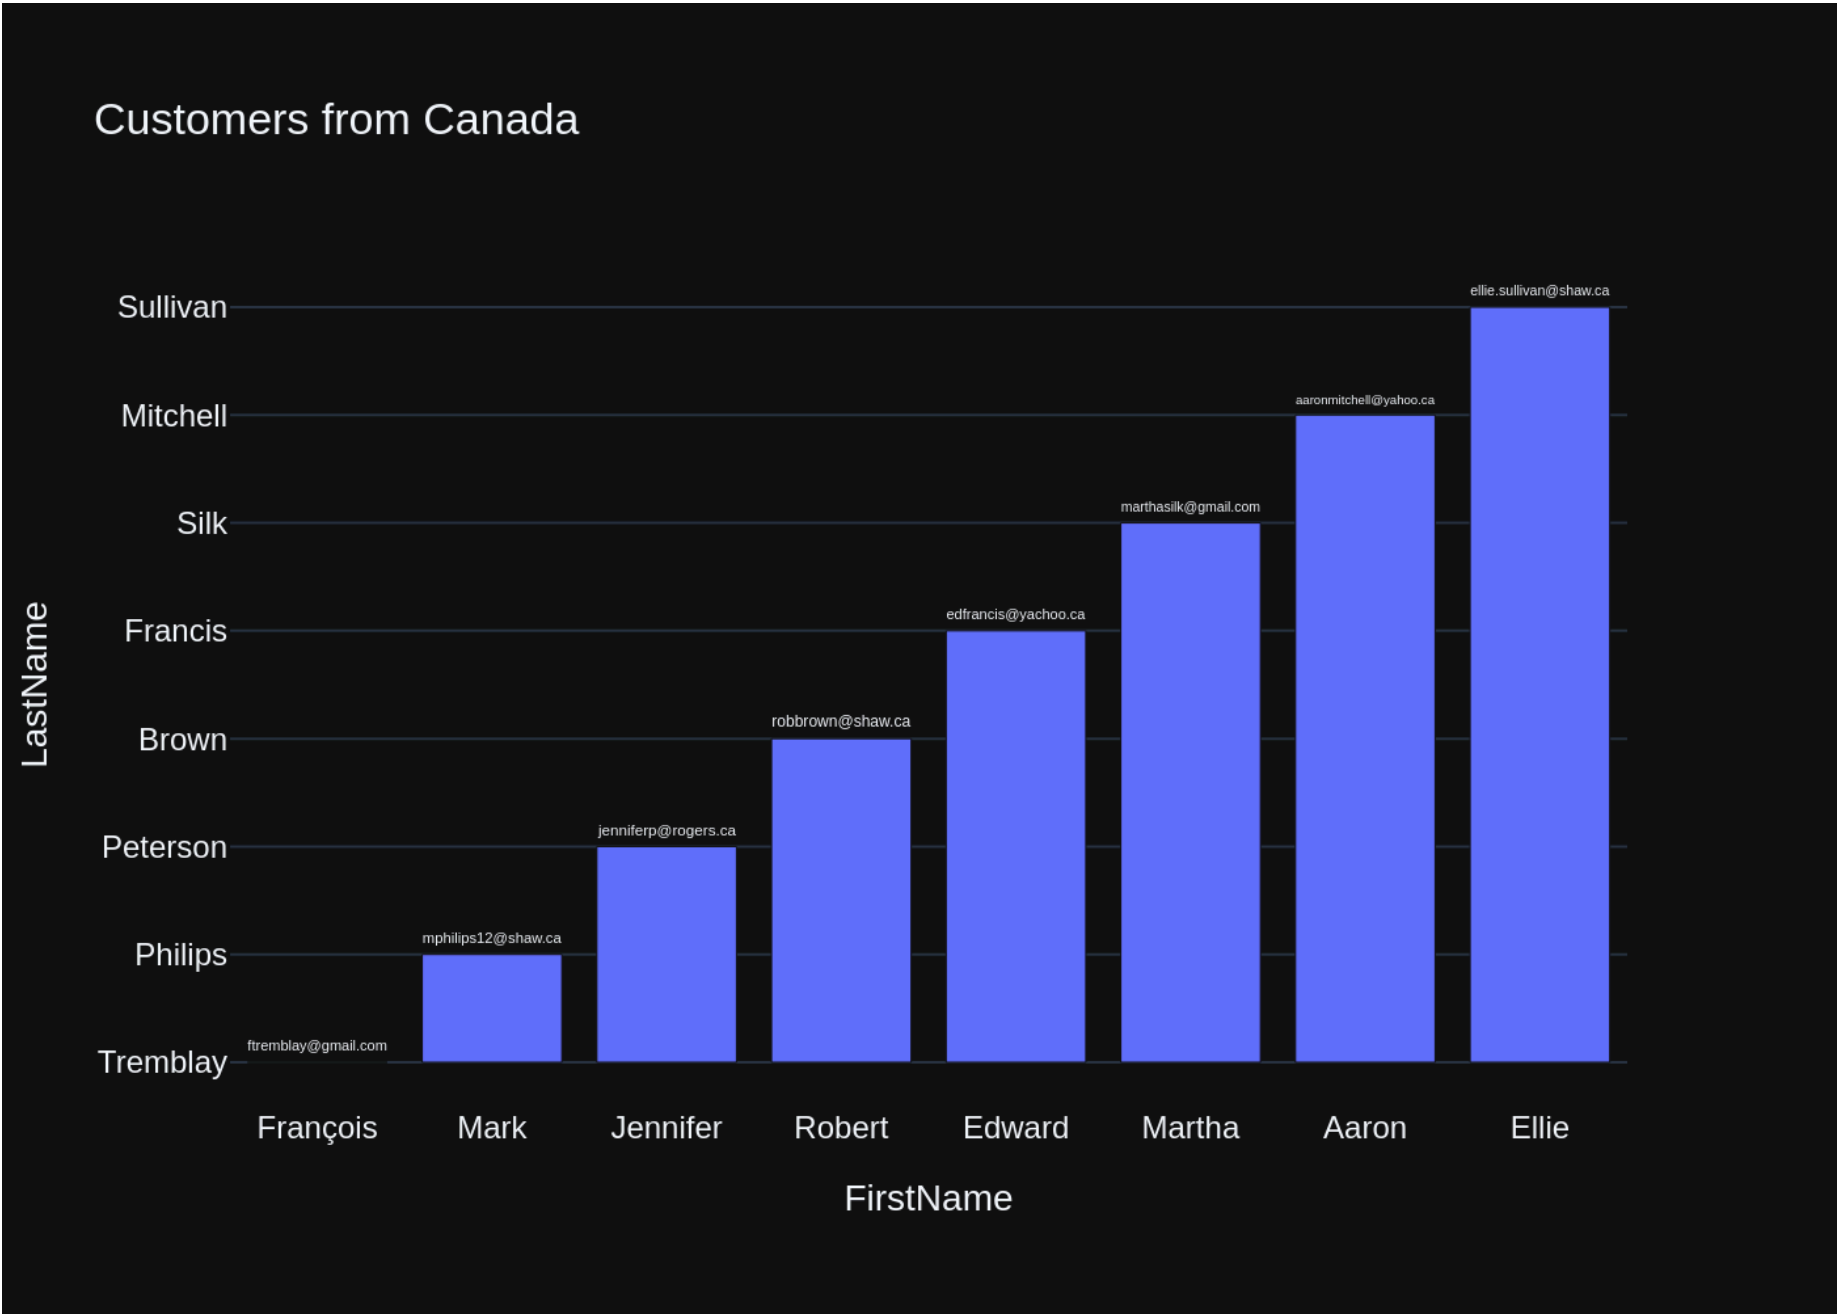
|   |        |      |                      |
|---|--------|------|----------------------|
| 5 | Martha | Silk | marthasilk@gmail.com |
|---|--------|------|----------------------|

|   |       |          |                        |
|---|-------|----------|------------------------|
| 6 | Aaron | Mitchell | aaronmitchell@yahoo.ca |
|---|-------|----------|------------------------|

|   |       |          |                        |
|---|-------|----------|------------------------|
| 7 | Ellie | Sullivan | ellie.sullivan@shaw.ca |
|---|-------|----------|------------------------|

Using model gpt-4o-mini for 190.5 tokens (approx)





```

Out[36]: ("SELECT FirstName, LastName, Email \nFROM customers \nWHERE Country = 'Canada';",
 FirstName LastName Email
0 François Tremblay ftremblay@gmail.com
1 Mark Philips mphilips12@shaw.ca
2 Jennifer Peterson jenniferp@rogers.ca
3 Robert Brown robbrown@shaw.ca
4 Edward Francis edfrancis@yachoo.ca
5 Martha Silk marthasilk@gmail.com
6 Aaron Mitchell aaronmitchell@yahoo.ca
7 Ellie Sullivan ellie.sullivan@shaw.ca,
Figure({
 'data': [{'alignmentgroup': 'True',
 'hovertemplate': 'FirstName=%{x}
LastName=%{y}
Email=%{text}<extra></extra>',
 'legendgroup': '',
 'marker': {'color': '#636efa', 'pattern': {'shape': ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'text': array(['ftremblay@gmail.com', 'mphilips12@shaw.ca', 'jenniferp@rogers.ca',
 'robbrown@shaw.ca', 'edfrancis@yachoo.ca', 'marthasilk@gmail.com',
 'aaronmitchell@yahoo.ca', 'ellie.sullivan@shaw.ca'], dtype=object),
 'textposition': 'outside',
 'texttemplate': '%{text}',
 'type': 'bar',
 'x': array(['François', 'Mark', 'Jennifer', 'Robert', 'Edward', 'Martha', 'Aaron',
 'Ellie'], dtype=object),
 'xaxis': 'x',
 'y': array(['Tremblay', 'Philips', 'Peterson', 'Brown', 'Francis', 'Silk',
 'Mitchell', 'Sullivan'], dtype=object),
 'yaxis': 'y'}],
 'layout': {'barmode': 'relative',
 'legend': {'tracegroupgap': 0},
 'template': '...',
 'title': {'text': 'Customers from Canada'},
 'xaxis': {'anchor': 'y', 'domain': [0.0, 1.0], 'title': {'text': 'FirstName'}},
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'LastName'}}
 })

```

```
In [37]: question = """
 Find the customer with the most invoices
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

SQL Prompt: [{'role': 'system', 'content': 'You are a SQLite expert. Please help to generate a SQL query to answer the question. Your response should ONLY be based on the given context and follow the response guidelines and format instructions. \n===Tables\nCREATE TABLE "invoices"\n\n InvoiceId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n CustomerId INTEGER NOT NULL,\n InvoiceDate DATETIME NOT NULL,\n BillingAddress NVARCHAR(70),\n BillingCity NVARCHAR(40),\n BillingState NVARCHAR(40),\n BillingCountry NVARCHAR(40),\n BillingPostalCode NVARCHAR(10),\n Total NUMERIC(10,2) NOT NULL,\n FOREIGN KEY (CustomerId) REFERENCES "customers" (CustomerId)\n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE INDEX IFK\_InvoiceCustomerId ON "invoices" (CustomerId)\n\nCREATE INDEX IFK\_InvoiceLineInvoiceId ON "invoice\_items" (InvoiceId)\n\nCREATE TABLE "invoice\_items"\n\n InvoiceLineId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n InvoiceId INTEGER NOT NULL,\n TrackId INTEGER NOT NULL,\n UnitPrice NUMERIC(10,2) NOT NULL,\n Quantity INTEGER NOT NULL,\n FOREIGN KEY (InvoiceId) REFERENCES "invoices" (InvoiceId)\n\nON DELETE NO ACTION ON UPDATE NO ACTION,\n FOREIGN KEY (TrackId) REFERENCES "tracks" (TrackId)\n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE INDEX IFK\_InvoiceLineTrackId ON "invoice\_items" (TrackId)\n\nCREATE TABLE "customers"\n\n CustomerId INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,\n FirstName NVARCHAR(40) NOT NULL,\n LastName NVARCHAR(20) NOT NULL,\n Company NVARCHAR(80),\n Address NVARCHAR(70),\n City NVARCHAR(40),\n State NVARCHAR(40),\n Country NVARCHAR(40),\n PostalCode NVARCHAR(10),\n Phone NVARCHAR(24),\n Fax NVARCHAR(24),\n Email NVARCHAR(60) NOT NULL,\n SupportRepId INTEGER,\n FOREIGN KEY (SupportRepId) REFERENCES "employees" (EmployeeId)\n\nON DELETE NO ACTION ON UPDATE NO ACTION\n\nCREATE INDEX IFK\_CustomerSupportRepId ON "customers" (SupportRepId)\n\n\n===Additional Context\n\n===Response Guidelines\n1. If the provided context is sufficient, please generate a valid SQL query without any explanations for the question.\n2. If the provided context is almost sufficient but requires knowledge of a specific string in a particular column, please generate an intermediate SQL query to find the distinct strings in that column. Prepend the query with a comment saying intermediate\_sql\n3. If the provided context is insufficient, please explain why it can't be generated.\n4. Please use the most relevant table(s).\n5. If the question has been asked and answered before, please repeat the answer exactly as it was given before.\n'}], [{'role': 'user', 'content': '\n Get the total number of invoices for each customer\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(\*) AS TotalInvoices\nFROM invoices\nGROUP BY CustomerId;'}, {'role': 'user', 'content': '\n List all invoices with a total exceeding \$10:\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': '\n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': 'SELECT \*, Total\nFROM invoices\nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': '\n Find the total number of invoices per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(\*) AS TotalInvoices\nFROM invoices\nGROUP BY BillingCountry;'}, {'role': 'user', 'content': '\n Get the average invoice total for each customer:\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, AVG(Total) AS AverageInvoiceTotal\nFROM invoices\nGROUP BY CustomerId;'}, {'role': 'user', 'content': '\n Find the top 5 most expensive tracks (based on unit price):\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM tracks ORDER BY UnitPrice DESC LIMIT 5;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': 'How many records are in table called customer'}, {'role': 'assistant', 'content': 'SELECT COUNT(\*) FROM customers;'}, {'role': 'user', 'content': '\n List all customers from Canada and their email addresses:\n'}, {'role': 'assistant', 'content': 'SELECT FirstName, LastName, Email\nFROM customers\nWHERE Country = 'Canada';"}, {'role': 'user', 'content': '\n List all employees and their reporting manager's name (if any):\n'}, {'role': 'assistant', 'content': 'SELECT e.EmployeeId, e.FirstName AS EmployeeFirstName, e.LastName AS EmployeeLastName, m.FirstName AS ManagerFirstName\nFROM employees e\nLEFT JOIN employees m ON e.SupportRepId = m.EmployeeId\nORDER BY e.EmployeeId\n'}]

```
rstName AS ManagerFirstName, m.LastName AS ManagerLastName\nFROM employees e\nLEFT JOIN employees m ON e.ReportsTo =
m.EmployeeId;'}, {'role': 'user', 'content': ' \n Find the customer with the most invoices \n'}]]
```

Using model gpt-4o-mini for 1006.5 tokens (approx)

LLM Response: SELECT CustomerId, COUNT(\*) AS InvoiceCount

FROM invoices

GROUP BY CustomerId

ORDER BY InvoiceCount DESC

LIMIT 1;

Extracted SQL: SELECT CustomerId, COUNT(\*) AS InvoiceCount

FROM invoices

GROUP BY CustomerId

ORDER BY InvoiceCount DESC

LIMIT 1;

SELECT CustomerId, COUNT(\*) AS InvoiceCount

FROM invoices

GROUP BY CustomerId

ORDER BY InvoiceCount DESC

LIMIT 1;

|   | CustomerId | InvoiceCount |
|---|------------|--------------|
| 0 | 1          | 7            |

Using model gpt-4o-mini for 192.75 tokens (approx)

Customer ID: 1

7

▲7

```
Out[37]: ('SELECT CustomerId, COUNT(*) AS InvoiceCount \nFROM invoices \nGROUP BY CustomerId \nORDER BY InvoiceCount DESC
\nLIMIT 1;',
 CustomerId InvoiceCount
 0 1 7,
 Figure({
 'data': [{'delta': {'reference': 0},
 'mode': 'number+delta',
 'title': {'text': 'Customer ID: 1'},
 'type': 'indicator',
 'value': 7}],
 'layout': {'template': '...'}
 })
```

In [ ]:

## Advanced SQL questions

```
In [38]: question = """
 Find the customer who bought the most albums in total quantity (across all invoices):
 """
 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

104/129



sponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': ' \nList all genres and the number of tracks in each genre:\n'}, {'role': 'assistant', 'content': 'SELECT genres.Name AS GenreName, COUNT(tracks.TrackId) AS TrackCount \nFROM genres \nLEFT JOIN tracks ON genres.GenreId = tracks.GenreId \nGROUP BY genres.GenreId;'}, {'role': 'user', 'content': ' \nFind the customer who bought the most albums in total quantity (across all invoices): \n'}]

Using model gpt-4o-mini for 1106.5 tokens (approx)

LLM Response: SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums  
FROM invoice\_items AS InvoiceItems  
JOIN invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId  
GROUP BY Invoice.CustomerId  
ORDER BY TotalAlbums DESC  
LIMIT 1;

Extracted SQL: SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums  
FROM invoice\_items AS InvoiceItems  
JOIN invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId  
GROUP BY Invoice.CustomerId  
ORDER BY TotalAlbums DESC  
LIMIT 1;

SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums  
FROM invoice\_items AS InvoiceItems  
JOIN invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId  
GROUP BY Invoice.CustomerId  
ORDER BY TotalAlbums DESC  
LIMIT 1;

| CustomerId | TotalAlbums |
|------------|-------------|
| 0          | 138         |

Using model gpt-4o-mini for 234.75 tokens (approx)

Customer ID: 1  
Total Albums Bought

38

```
Out[38]: ('SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums \nFROM invoice_items AS InvoiceItems \nJOIN
invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId \nGROUP BY Invoice.CustomerId \nORDER BY TotalAl
bums DESC \nLIMIT 1;',
CustomerId TotalAlbums
0 1 38,
Figure({
 'data': [{'mode': 'number',
 'title': {'text': 'Customer ID: 1
Total Albums Bought'},
 'type': 'indicator',
 'value': 38}],
 'layout': {'template': '...'}
}))
```

```
In [39]: question = """
 Find the top 5 customer who bought the most albums in total quantity (across all invoices):
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

108/129

```

artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName
\nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': ' \n Find all
invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': "SELECT *, Total \nFROM inv
oices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n Find the total number of invoices
per country:\n'}, {'role': 'assistant', 'content': 'SELECT BillingCountry, COUNT(*) AS TotalInvoices \nFROM invoices
\nGROUP BY BillingCountry;'}, {'role': 'user', 'content': ' \n Find the top 5 customer who bought the most albu
ms in total quantity (across all invoices):\n'}]

```

Using model gpt-4o-mini for 1136.25 tokens (approx)

```

LLM Response: SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums
FROM invoice_items AS InvoiceItems
JOIN invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId
GROUP BY Invoice.CustomerId
ORDER BY TotalAlbums DESC
LIMIT 5;

```

```

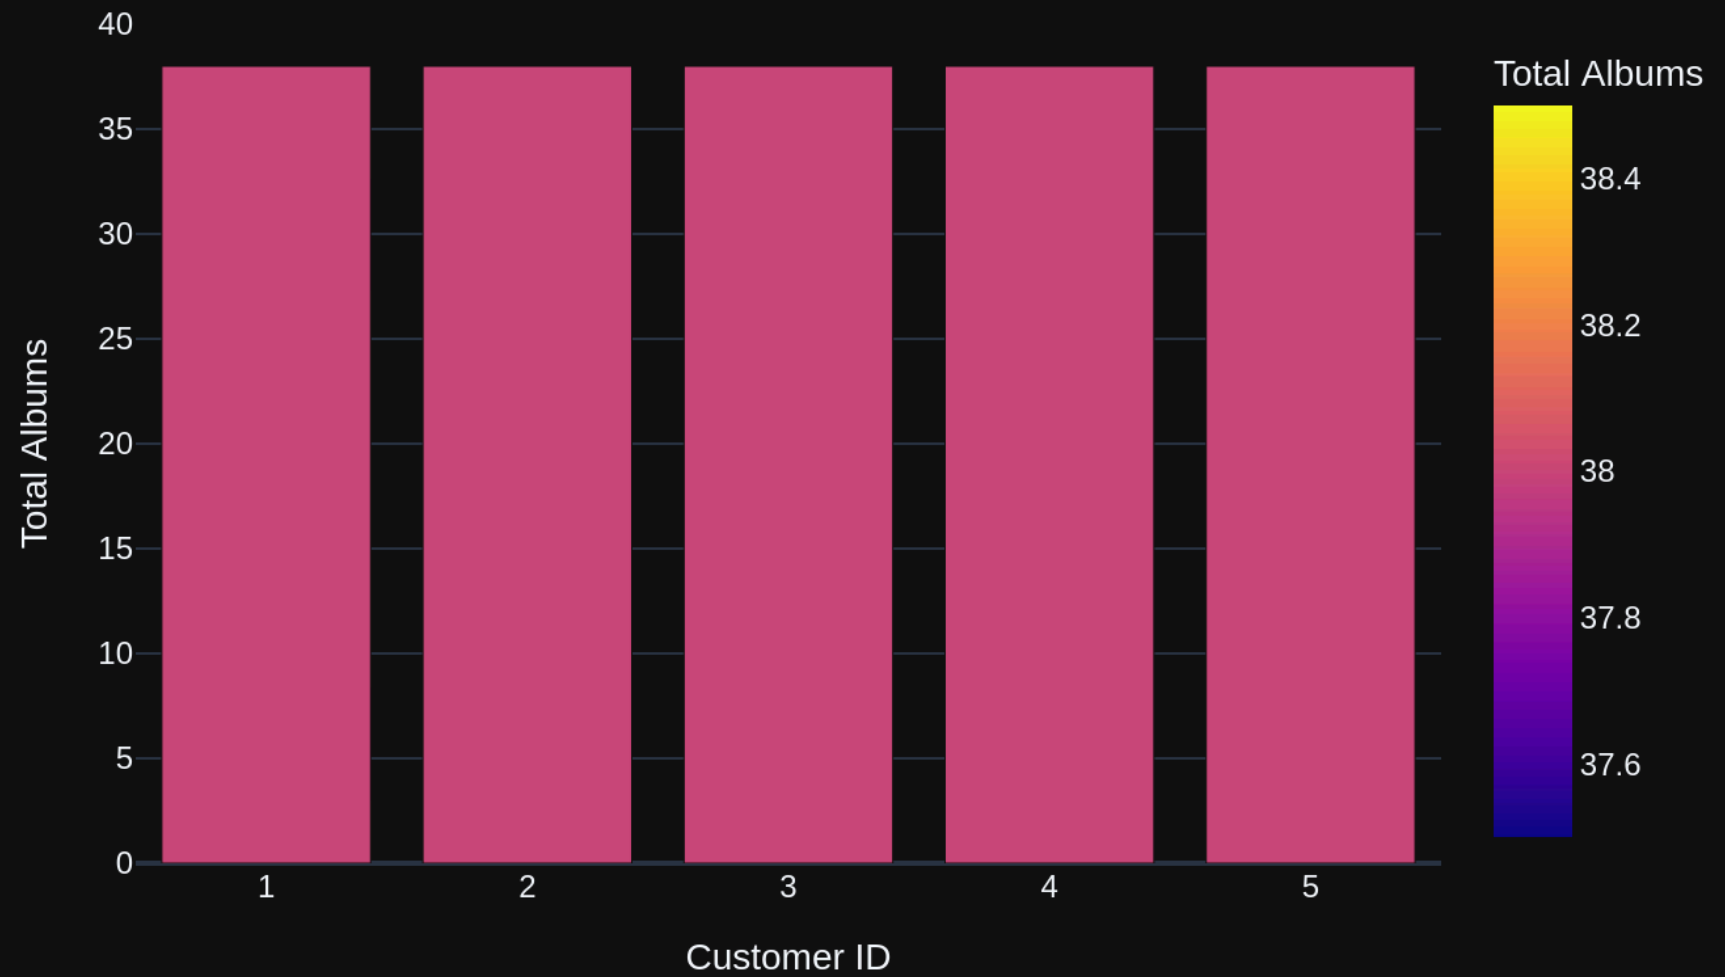
Extracted SQL: SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums
FROM invoice_items AS InvoiceItems
JOIN invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId
GROUP BY Invoice.CustomerId
ORDER BY TotalAlbums DESC
LIMIT 5;
SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums
FROM invoice_items AS InvoiceItems
JOIN invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId
GROUP BY Invoice.CustomerId
ORDER BY TotalAlbums DESC
LIMIT 5;

```

|   | CustomerId | TotalAlbums |
|---|------------|-------------|
| 0 | 1          | 38          |
| 1 | 2          | 38          |
| 2 | 3          | 38          |
| 3 | 4          | 38          |
| 4 | 5          | 38          |

Using model gpt-4o-mini for 236.0 tokens (approx)

## Top 5 Customers by Total Albums Purchased



```
Out[39]: ('SELECT Invoice.CustomerId, SUM(InvoiceItems.Quantity) AS TotalAlbums \nFROM invoice_items AS InvoiceItems \nJOIN
invoices AS Invoice ON InvoiceItems.InvoiceId = Invoice.InvoiceId \nGROUP BY Invoice.CustomerId \nORDER BY TotalAl
bums DESC \nLIMIT 5;',
```

|   | CustomerId | TotalAlbums |
|---|------------|-------------|
| 0 | 1          | 38          |
| 1 | 2          | 38          |
| 2 | 3          | 38          |
| 3 | 4          | 38          |
| 4 | 5          | 38,         |

```
Figure({
 'data': [{ 'alignmentgroup': 'True',
 'hovernment': 'Customer ID=%{x}
Total Albums=%{marker.color}<extra></extra>',
 'legendgroup': '',
 'marker': { 'color': array([38, 38, 38, 38, 38]), 'coloraxis': 'coloraxis', 'pattern': { 'shape':
 ''}},
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array([1, 2, 3, 4, 5]),
 'xaxis': 'x',
 'y': array([38, 38, 38, 38, 38]),
 'yaxis': 'y'}],
 'layout': { 'barmode': 'relative',
 'coloraxis': { 'colorbar': { 'title': { 'text': 'Total Albums' } },
 'colorscale': [[0.0, '#0d0887'], [0.11111111111111111,
 '#46039f'], [0.22222222222222222,
 '#7201a8'], [0.33333333333333333,
 '#9c179e'], [0.44444444444444444,
 '#bd3786'], [0.55555555555555556,
 '#d8576b'], [0.66666666666666666,
 '#ed7953'], [0.77777777777777778,
 '#fb9f3a'], [0.88888888888888888,
 '#fdca26'], [1.0, '#f0f921']]],
 'legend': { 'tracegroupgap': 0 },
 'template': '...',
 'title': { 'text': 'Top 5 Customers by Total Albums Purchased' },
 'xaxis': { 'anchor': 'y', 'domain': [0.0, 1.0], 'title': { 'text': 'Customer ID' } },
```

```
 'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Total Albums'}}})
```

```
In [40]: question = """
 Find the top 3 customers who spent the most money overall:
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1



file:///home/gongai/Downloads/openai-gpt-4o-mini-chromadb-sqlite-test-1.html

```
\n Get the total number of invoices for each customer\n'}, {'role': 'assistant', 'content': 'SELECT CustomerId, COUNT(*) AS TotalInvoices \nFROM invoices \nGROUP BY CustomerId;'}, {'role': 'user', 'content': 'How many customers are there'}, {'role': 'assistant', 'content': 'SELECT COUNT(*) FROM customers;'}, {'role': 'user', 'content': ' \nList all invoices with a total exceeding $10:\n'}, {'role': 'assistant', 'content': 'SELECT * FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': ' \n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': 'SELECT *, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n Find the top 3 customers who spent the most money overall:\n'}]
```

Using model gpt-4o-mini for 1146.75 tokens (approx)

LLM Response: SELECT CustomerId, SUM(Total) AS TotalSpent

FROM invoices

GROUP BY CustomerId

ORDER BY TotalSpent DESC

LIMIT 3;

Extracted SQL: SELECT CustomerId, SUM(Total) AS TotalSpent

FROM invoices

GROUP BY CustomerId

ORDER BY TotalSpent DESC

LIMIT 3;

SELECT CustomerId, SUM(Total) AS TotalSpent

FROM invoices

GROUP BY CustomerId


ORDER BY TotalSpent DESC

LIMIT 3;

|   | CustomerId | TotalSpent |
|---|------------|------------|
| 0 | 6          | 49.62      |
| 1 | 26         | 47.62      |
| 2 | 57         | 46.62      |

Using model gpt-4o-mini for 196.5 tokens (approx)

### Top 3 Customers by Total Spending



```
Out[40]: ('SELECT CustomerId, SUM(Total) AS TotalSpent \nFROM invoices \nGROUP BY CustomerId \nORDER BY TotalSpent DESC \nL
IMIT 3;',
 CustomerId TotalSpent
0 6 49.62
1 26 47.62
2 57 46.62,
 Figure({
 'data': [{'type': 'bar', 'x': array(['6', '26', '57'], dtype=object), 'y': array([49.62, 47.62, 46.62])}],
 'layout': {'template': '...',
 'title': {'text': 'Top 3 Customers by Total Spending'},
 'xaxis': {'title': {'text': 'Customer ID'}},
 'yaxis': {'title': {'text': 'Total Spent'}}}
 })
```

```
In [41]: question = """
 Get all playlists containing at least 10 tracks and the total duration of those tracks:
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

117/129

t all albums and their corresponding artist names \n'}, {'role': 'assistant', 'content': 'SELECT albums.Title AS AlbumTitle, artists.Name AS ArtistName \nFROM albums \nJOIN artists ON albums.ArtistId = artists.ArtistId;'}, {'role': 'user', 'content': ' \n Find all invoices since 2010 and the total amount invoiced:\n'}, {'role': 'assistant', 'content': "SELECT \*, Total \nFROM invoices \nWHERE InvoiceDate >= '2010-01-01';"}, {'role': 'user', 'content': ' \n List all invoices with a total exceeding \$10:\n'}, {'role': 'assistant', 'content': 'SELECT \* FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': 'Show me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite\_master WHERE type='table';"}, {'role': 'user', 'content': ' \n Get all playlists containing at least 10 tracks and the total duration of those tracks:\n'}]

Using model gpt-4o-mini for 1174.5 tokens (approx)

LLM Response: SELECT playlists.PlaylistId, playlists.Name, COUNT(playlist\_track.TrackId) AS TrackCount, SUM(tracks.Milliseconds) AS TotalDuration

FROM playlists

JOIN playlist\_track ON playlists.PlaylistId = playlist\_track.PlaylistId

JOIN tracks ON playlist\_track.TrackId = tracks.TrackId

GROUP BY playlists.PlaylistId

HAVING TrackCount >= 10;

Extracted SQL: SELECT playlists.PlaylistId, playlists.Name, COUNT(playlist\_track.TrackId) AS TrackCount, SUM(tracks.Milliseconds) AS TotalDuration

FROM playlists

JOIN playlist\_track ON playlists.PlaylistId = playlist\_track.PlaylistId

JOIN tracks ON playlist\_track.TrackId = tracks.TrackId

GROUP BY playlists.PlaylistId

HAVING TrackCount >= 10;

SELECT playlists.PlaylistId, playlists.Name, COUNT(playlist\_track.TrackId) AS TrackCount, SUM(tracks.Milliseconds) AS TotalDuration

FROM playlists

JOIN playlist\_track ON playlists.PlaylistId = playlist\_track.PlaylistId


JOIN tracks ON playlist\_track.TrackId = tracks.TrackId

GROUP BY playlists.PlaylistId

HAVING TrackCount >= 10;

|   | PlaylistId | Name                       | TrackCount | TotalDuration |
|---|------------|----------------------------|------------|---------------|
| 0 | 1          | Music                      | 3290       | 877683083     |
| 1 | 3          | TV Shows                   | 213        | 501094957     |
| 2 | 5          | 90's Music                 | 1477       | 398705153     |
| 3 | 8          | Music                      | 3290       | 877683083     |
| 4 | 10         | TV Shows                   | 213        | 501094957     |
| 5 | 11         | Brazilian Music            | 39         | 9486559       |
| 6 | 12         | Classical                  | 75         | 21770592      |
| 7 | 13         | Classical 101 - Deep Cuts  | 25         | 6755730       |
| 8 | 14         | Classical 101 - Next Steps | 25         | 7575051       |
| 9 | 15         | Classical 101 - The Basics | 25         | 7439811       |

|                                                    |    |                     |    |         |
|----------------------------------------------------|----|---------------------|----|---------|
| 10                                                 | 16 | Grunge              | 15 | 4122018 |
| 11                                                 | 17 | Heavy Metal Classic | 26 | 8206312 |
| Using model gpt-4o-mini for 271.25 tokens (approx) |    |                     |    |         |





```
Out[41]: ('SELECT playlists.PlaylistId, playlists.Name, COUNT(playlist_track.TrackId) AS TrackCount, SUM(tracks.Millisecond
s) AS TotalDuration \nFROM playlists \nJOIN playlist_track ON playlists.PlaylistId = playlist_track.PlaylistId \nJ
OIN tracks ON playlist_track.TrackId = tracks.TrackId \nGROUP BY playlists.PlaylistId \nHAVING TrackCount >= 10;',
```

|    | PlaylistId | Name                       | TrackCount | TotalDuration |
|----|------------|----------------------------|------------|---------------|
| 0  | 1          | Music                      | 3290       | 877683083     |
| 1  | 3          | TV Shows                   | 213        | 501094957     |
| 2  | 5          | 90's Music                 | 1477       | 398705153     |
| 3  | 8          | Music                      | 3290       | 877683083     |
| 4  | 10         | TV Shows                   | 213        | 501094957     |
| 5  | 11         | Brazilian Music            | 39         | 9486559       |
| 6  | 12         | Classical                  | 75         | 21770592      |
| 7  | 13         | Classical 101 - Deep Cuts  | 25         | 6755730       |
| 8  | 14         | Classical 101 - Next Steps | 25         | 7575051       |
| 9  | 15         | Classical 101 - The Basics | 25         | 7439811       |
| 10 | 16         | Grunge                     | 15         | 4122018       |
| 11 | 17         | Heavy Metal Classic        | 26         | 8206312,      |

```
Figure({
 'data': [{ 'alignmentgroup': 'True',
 'hovernment': 'Playlist Name={x}
Total Duration (Milliseconds)={y}<extra></extra>',
 'legendgroup': '',
 'marker': { 'color': '#636efa', 'pattern': { 'shape': '' } },
 'name': '',
 'offsetgroup': '',
 'orientation': 'v',
 'showlegend': False,
 'textposition': 'auto',
 'type': 'bar',
 'x': array(['Music', 'TV Shows', '90's Music', 'Music', 'TV Shows',
 'Brazilian Music', 'Classical', 'Classical 101 - Deep Cuts',
 'Classical 101 - Next Steps', 'Classical 101 - The Basics', 'Grunge',
 'Heavy Metal Classic'], dtype=object),
 'xaxis': 'x',
 'y': array([877683083, 501094957, 398705153, 877683083, 501094957, 9486559,
 21770592, 6755730, 7575051, 7439811, 4122018, 8206312]),
 'yaxis': 'y'}],
 'layout': { 'barmode': 'relative',
 'legend': { 'tracegroupgap': 0 },
 'template': '...',
 'title': { 'text': 'Total Duration of Playlists with at least 10 Tracks' },
 'xaxis': { 'anchor': 'y', 'domain': [0.0, 1.0], 'title': { 'text': 'Playlist Name' } },
```

```
'yaxis': {'anchor': 'x', 'domain': [0.0, 1.0], 'title': {'text': 'Total Duration (Millisecond
s)}}}}
}))
```

```
In [42]: question = """
 Identify artists who have albums with tracks appearing in multiple genres:
 """

 vn.ask(question=question)
```

Number of requested results 10 is greater than number of elements in index 1, updating n\_results = 1

123/129


```
merId \nORDER BY TotalAlbums DESC \nLIMIT 5;'}, {'role': 'user', 'content': ' \n Find the top 5 most expensive t
racks (based on unit price):\n'}, {'role': 'assistant', 'content': 'SELECT * FROM tracks ORDER BY UnitPrice DESC LIM
IT 5;'}, {'role': 'user', 'content': ' \n Find all tracks with a name containing "What" (case-insensitive)\n'},
{'role': 'assistant', 'content': "SELECT * FROM tracks WHERE Name LIKE '%What%';"}, {'role': 'user', 'content': 'Sho
w me a list of tables in the SQLite database'}, {'role': 'assistant', 'content': "SELECT name FROM sqlite_master WHE
RE type='table';"}, {'role': 'user', 'content': ' \n List all invoices with a total exceeding $10:\n'}, {'role':
'assistant', 'content': 'SELECT * FROM invoices WHERE Total > 10;'}, {'role': 'user', 'content': ' \n Identify
artists who have albums with tracks appearing in multiple genres:\n'}]
```

Using model gpt-4o-mini for 1175.5 tokens (approx)

```
LLM Response: SELECT DISTINCT artists.Name AS ArtistName
FROM artists
JOIN albums ON artists.ArtistId = albums.ArtistId
JOIN tracks ON albums.AlbumId = tracks.AlbumId
GROUP BY artists.ArtistId
HAVING COUNT(DISTINCT tracks.GenreId) > 1;
Extracted SQL: SELECT DISTINCT artists.Name AS ArtistName
FROM artists
JOIN albums ON artists.ArtistId = albums.ArtistId
JOIN tracks ON albums.AlbumId = tracks.AlbumId
GROUP BY artists.ArtistId
HAVING COUNT(DISTINCT tracks.GenreId) > 1;
SELECT DISTINCT artists.Name AS ArtistName
FROM artists
JOIN albums ON artists.ArtistId = albums.ArtistId
JOIN tracks ON albums.AlbumId = tracks.AlbumId
GROUP BY artists.ArtistId
HAVING COUNT(DISTINCT tracks.GenreId) > 1;
```

|    | ArtistName           |
|----|----------------------|
| 0  | Antônio Carlos Jobim |
| 1  | Audioslave           |
| 2  | Various Artists      |
| 3  | Gilberto Gil         |
| 4  | Eric Clapton         |
| 5  | Faith No More        |
| 6  | Foo Fighters         |
| 7  | Guns N' Roses        |
| 8  | Iron Maiden          |
| 9  | Jamiroquai           |
| 10 | Lenny Kravitz        |
| 11 | Ozzy Osbourne        |
| 12 | Pearl Jam            |

13                   R.E.M.  
14 Red Hot Chili Peppers  
15   Battlestar Galactica  
16                   Heroes  
17                   Lost  
18                   U2  
19           The Office  
20           Amy Winehouse  
Using model gpt-4o-mini for 222.5 tokens (approx)



```
Out[42]: ('SELECT DISTINCT artists.Name AS ArtistName \nFROM artists \nJOIN albums ON artists.ArtistId = albums.ArtistId \n
JOIN tracks ON albums.AlbumId = tracks.AlbumId \nGROUP BY artists.ArtistId \nHAVING COUNT(DISTINCT tracks.GenreId)
> 1;',
```

```

 ArtistName
0 Antônio Carlos Jobim
1 Audioslave
2 Various Artists
3 Gilberto Gil
4 Eric Clapton
5 Faith No More
6 Foo Fighters
7 Guns N' Roses
8 Iron Maiden
9 Jamiroquai
10 Lenny Kravitz
11 Ozzy Osbourne
12 Pearl Jam
13 R.E.M.
14 Red Hot Chili Peppers
15 Battlestar Galactica
16 Heroes
17 Lost
18 U2
19 The Office
20 Amy Winehouse,
Figure({
 'data': [{'type': 'bar',
 'x': array(['Antônio Carlos Jobim', 'Audioslave', 'Various Artists', 'Gilberto Gil',
 'Eric Clapton', 'Faith No More', 'Foo Fighters', 'Guns N' Roses',
 'Iron Maiden', 'Jamiroquai', 'Lenny Kravitz', 'Ozzy Osbourne',
 'Pearl Jam', 'R.E.M.', 'Red Hot Chili Peppers', 'Battlestar Galactica',
 'Heroes', 'Lost', 'U2', 'The Office', 'Amy Winehouse'], dtype=object),
 'y': [1, 1]}],
 'layout': {'template': '...',
 'title': {'text': 'Artists with Albums in Multiple Genres'},
 'xaxis': {'title': {'text': 'Artist Name'}},
 'yaxis': {'showgrid': False, 'title': {'text': 'Count'}, 'zeroline': False}}
)))
```

## Check completion time

```
In [43]: ts_stop = time()

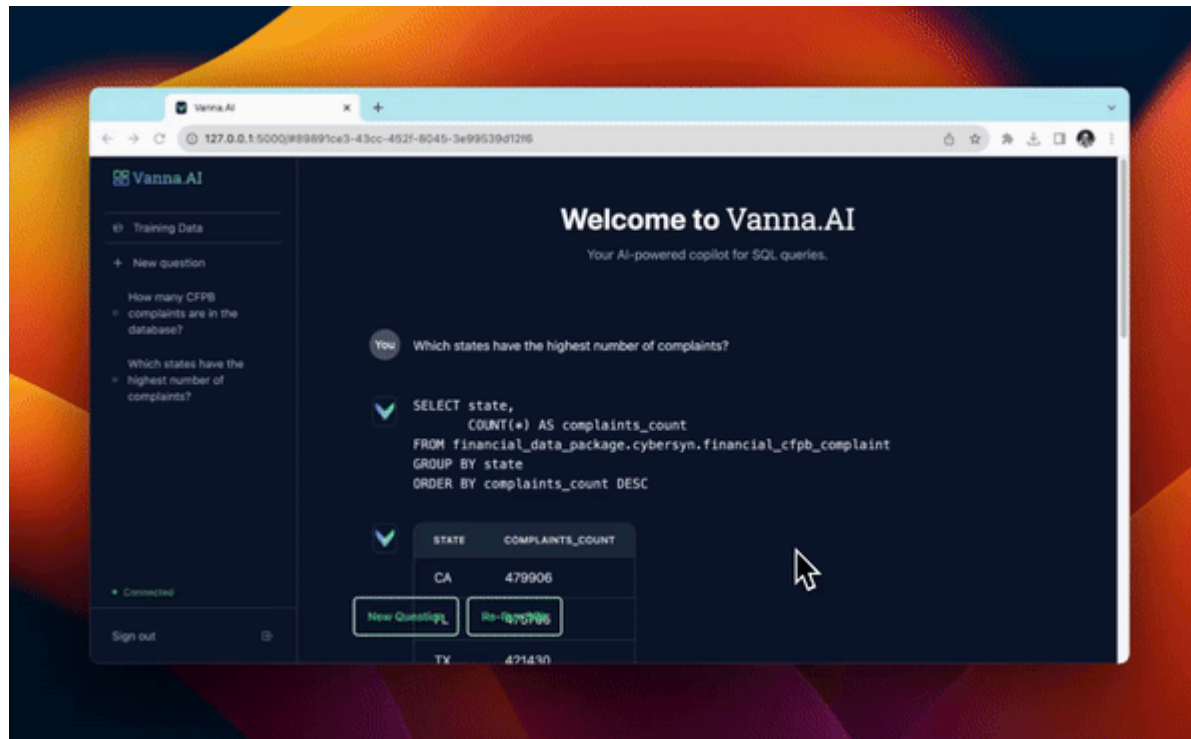
elapsed_time = ts_stop - ts_start
print(f"test running on '{hostname}' with '{model_name}' LLM took : {elapsed_time:.2f} sec")
```

test running on 'ducklover1' with 'gpt-4o-mini' LLM took : 73.75 sec

```
In [44]: from datetime import datetime
print(datetime.now())
```

2024-07-21 20:59:35.225156

## Launch the User Interface



```
from vanna.flask import VannaFlaskApp app = VannaFlaskApp(vn) app.run()
```



## Next Steps

Using Vanna via Jupyter notebooks is great for getting started but check out additional customizable interfaces like the

- [Streamlit app](#)
- [Flask app](#)
- [Slackbot](#)